

# Exploration, Grouping, and Prediction of Genre in Spotify Songs

**Team Spotiflies: Aubrey Clark, Kennedy Kang, Aaron Kwan, Joanna Liu, Ethan Mai, Aster Phan**

- [GitHub Link](#)

## I

In this project, we predict song genre using a [massive dataset](#) of **over 140 thousand songs** from the music streaming platform Spotify. The data set had information about each song's popularity, genre, and other musical qualities. We attempted to classify a song into one of 10 genres by using up to 12 different predictor variables.

- Relevant Predictor Variables: popularity , duration\_ms , explicit , danceability , energy , loudness , speechiness , acousticness , instrumentalness , liveness , valence , tempo
- Genres: classical , electronic , folk , hip-hop , jazz , metal , misc , pop , rock , world

## II

The main feature we are interested in with this dataset is **song genre**, in particular how it relates to other features. We want to test if the musical qualities of a song are strong predictors of its genre. After cleaning, our dataset contained 114 different genres, ranging from emo to classical to industrial and more. We are interested in investigating things such as:

- The qualities of songs that are correlated to genre.
- If certain genres are more similar to each other than others.
- If we can predict a song's genre through its other qualities.

Understanding these relationships is crucial for several reasons. First, accurate genre prediction can *enhance recommendation systems*, providing users with a more personalized and enjoyable listening experience. For artists and producers, insights from our analysis can *guide the creative process*, helping them understand trends and preferences within different genres. Finally, we contribute to the broader field of musicology by providing *data-driven insights* into the characteristics that are different in musical genres.

## III

We tried numerous different methodologies, including **linear regression**, **logistic regression**, **decision trees**, **PCA/clustering**, and **neural networks**. Our key methodologies were **decision trees** and **random forest**, which worked most effectively to predict a song's genre through its other qualities.

With decision trees, we were able to obtain a respectable accuracy compared to our other methods. However, due to the depth of the decision trees, these results were not easily interpretable. Finally, after using an ensemble of trees in a **random forest**, we achieved predictions that were significantly more robust and accurate than the previous methods.

**Neural networks** also appeared to be a viable alternative, however overall accuracy metrics paled in comparison to decision trees. However, recall scores for classical music and metal were better than they were for **random forests**.

## IV

Using a train-test split of 60:20:20 and an ensemble of 200 trees, each with a maximum tree depth of 20 (both parameters learned using brute-force cross validation), we were able to achieve a total **Accuracy** of `0.56` . We generated a **Classification Report** showing the `precision` , `recall` , and `F1 score` for each genre, which helped us understand the performance of our model across different genres. In addition, we generated a **Confusion Matrix**, which provided a visual representation of the true positive, true negative, false positive, and false negative prediction rates, allowing us to see what errors our model was making.

Finally, we used **Feature & Permutation Importance** to identify which predictor variables had the most significant impact on genre classification. In particular, we found that `popularity` , `acousticness` , `instrumentalness` , and `danceability` were the most important to predict genre. This makes sense because some genres are more appealing to wider audiences and these features could logically seem to be able to predict between genres. Interestingly, `mode`, `time signature`, and `key` had little or negative permutation importances, indicating they are not important to predicting track genre.

While our Random Forest model provided robust and accurate predictions, it also had some limitations. One significant limitation was the model's tendency to overfit (the training accuracy was over 90%), especially with a large number of trees and depth. We tried to combat this by limiting the *maximum tree depth*, which resulted in lower accuracy. In addition, the computational resources required for training and tuning the model were substantial (each run took dozens of minutes) which could be a constraint for larger datasets or more extensive hyperparameter optimization.

All in all, our maximum accuracy was still well below `0.7` , which may be a limitation of our dataset itself; this dataset may not capture all the nuances and variations between genres. Additionally, the subjective nature of genre classification means that there can be overlap and ambiguity between genres, making it challenging to reach high accuracy. Future work could involve refining the genre labels, incorporating more diverse and representative data, and experimenting with different model architectures and techniques to improve performance. Despite these limitations, our project demonstrates the potential of machine learning techniques in predicting song genres and provides a foundation for further exploration and improvement.

## V

To run the code in **project\_code.ipynb**, ensure that Jupyter Notebook or Jupyter Lab is installed, along with the necessary dependencies. First, clone this repository to the local machine and navigate to the project directory. Launch Jupyter Notebook by running `jupyter notebook` in your terminal, open **project\_code.ipynb**, and execute the cells sequentially. Alternatively, use the `Jupyter Notebook` extension and use "Run all cells."

## Appendix

### I

Boxplots were used to identify the spread of numerical data. Histograms of each feature were created to visualize the frequency of values for the data. By doing this, we determined the skewedness of each feature. A heatmap was used to determine correlations between features. For features that were more correlated, we created scatterplots and used linear and quadratic regressions to find preliminary relationships. We also created histograms of categorical data to see the frequency of each value for the categorical data. Using domain knowledge, we grouped the various genres into 10 overarching genres. Bar graphs were used to visualize mean values for each numerical feature for the overarching genres. A random seed was used to split the data into 60% for training, 20% for validation, and 20% for testing.

### II

We realized that many similar genres were classified into different groups, so we grouped the given genres into 10 overarching genres. These grouping are shown below. We removed the "Unnamed: 0" column, any rows with NaN values, and any songs with a duration of less than 30 seconds. We also removed songs with a tempo or time signature of 0, as these values do not make sense for songs.

```
{
  'pop': ['cantopop', 'j-pop', 'j-idol', 'k-pop', 'mandopop', 'pop', 'indie-pop',
'power-pop', 'pop-film', 'synth-pop'],
  'rock': ['alt-rock', 'alternative', 'hard-rock', 'indie', 'punk', 'j-rock', 'punk-
rock', 'psych-rock', 'rock', 'rock-n-roll', 'grunge', 'emo', 'rockabilly', 'guitar'],
  'metal': ['black-metal', 'death-metal', 'heavy-metal', 'metal', 'metalcore',
'grindcore'],
  'electronic': ['edm', 'electro', 'electronic', 'house', 'garage', 'j-dance',
'hardcore', 'hardstyle', 'industrial', 'techno', 'trance', 'dubstep', 'idm', 'minimal-
techno', 'progressive-house', 'chicago-house', 'deep-house', 'detroit-techno',
'disco', 'drum-and-bass', 'dub', 'club', 'dance', 'dancehall'],
  'hip-hop': ['hip-hop', 'rap', 'r-n-b', 'breakbeat', 'trip-hop'],
  'jazz': ['jazz', 'blues', 'soul', 'funk', 'ska', 'gospel'],
  'classical': ['classical', 'opera', 'piano'],
  'world': ['afrobeat', 'brazil', 'sertanejo', 'british', 'latin', 'latino',
'samba', 'salsa', 'reggae', 'reggaeton', 'tango', 'world-music', 'indian', 'iranian',
'turkish', 'malay', 'mpb', 'pagode', 'forro', 'french', 'german', 'spanish',
'swedish'],
  'folk': ['folk', 'bluegrass', 'country', 'singer-songwriter', 'songwriter',
'honky-tonk'],
  'misc': ['acoustic', 'ambient', 'anime', 'children', 'chill', 'comedy', 'disney',
'happy', 'party', 'study', 'sleep', 'show-tunes', 'new-age', 'kids', 'goth', 'groove',
'romance', 'sad']
}
```

### III

Popularity, duration, danceability, energy, loudness, speechiness, acousticness, instrumentality, liveness, valence, and tempo were used in both linear and quadratic regression. We used a sequential feature selector, forward selection, to determine the most influential features. We used  $R^2$  as the scoring metric for cross-validation to identify optimal features. Because we used this metric, we found that regularization was not necessary. Overall, we determined that the importance of certain features varied by genre group. For

example, energy was more prominent for rock and electronic genres, acousticness and liveness were more prominent for folk and classical genres, and danceability was more prominent for pop and hip-hop genres.

#### **IV**

For logistic regression, the track-genre column was encoded using a label encoder to map genres to numerical labels. Non-numeric columns were converted via one-hot encoding, and data was standardized. We calculated accuracy, recall, and F1 score and plotted ROC curves and found the AUC to determine the model's ability to distinguish between classes. The default solver lbfgs used L2 regularization, so we did use regularization. However, removing the regularization did not affect the metrics.

We found that there were different key features for each genre. Energy contributed largely to rock, metal, and electronic music. Danceability contributed largely to pop and hip-hop identification. Acousticness contributed largely to classical and folk music. Loudness contributed largely to metal music.

#### **V**

We used KNN, a decision tree, and a random forest to try to classify genres of our data. For all of these, we one-hot-encoded all categorical data. While this isn't necessary for the tree models, it is for KNN, since distance between two points for categorical data does not make sense if the data is numeric.

Out of these three, the random forests had the best accuracy on the test set at 0.56, and was the best method by far for classifying the data. As expected, KNN was inefficient in analyzing our data due to the dataset's high dimensionality, and the presence of a lot of categorical data where distance makes much less sense. Decision trees make sense as a good model for the data, since there is a lot of overlap and the decision tree can make minute distinctions at each split. And random forests, as an ensemble version of decision trees, make sense as being more accurate.

#### **VI**

We applied one-hot encoding and scaled and mean-centered the data before doing PCA. The scree-plot results were unclear as to how many clusters should be used, but the top 4 principal components accounted for about 30% of the variability. Plotting the principal components showed that there was not clear separation of clusters due to heavy overlap.

After applying logistic regression to the PC-transformed dataset, we still produced relatively low accuracy rate of about 30%; the silhouette scores were also low.

We tried applying a Gaussian Mixture model on the PC-transformed dataset, which resulted in visually similar plots, but very low silhouette score. Even reducing the dimensionality of the GMM resulted in lower silhouette scores. We concluded that PCA transforming the data would not really help with determining genre. Since we did not have many features to begin with, we also believed that feature reduction was not a priority.

We also applied K-means clustering, agglomerative (hierarchical) clustering, and gaussian mixture models to the non-PCA-transformed data. All of these had terrible rand index with genre and <0.25 silhouette scores. Thus, we concluded that clustering would not help us with our problem.

#### **VII**

We used the dataset without string features, and one-hot-encoded the categorical features.

We went through a number of models, but ReLU, softmax, and dropout all significantly decreased the accuracy, so we ended up with a simpler model. Our model uses 2 hidden layers of size 64 using sigmoid activation, and log softmax output activation.

We trained using the Adam optimizer using cross entropy loss, and reported accuracy metrics at the end. We also included weights to account for class imbalance. We learned an initial learning rate via pytorch's `lr_find` function, and used the StepLR learning rate scheduler which reduced learning rate by 90% every 10 epochs. Our final accuracy was about 20%.

However, looking at accuracy metrics, recall for some classes like classical and metal was really high, indicating that if the model predicts that a song is classical or metal, then it is likely that it is right.

Interestingly, the model never predicted pop or hip-hop. This might be due to class imbalance, though pop is the most populated class and hip-hop the least. It might also be that the class weights undercompensated for hip-hop and overcompensated for pop. However, this was very common for many runs using many different hyperparameters.

## VIII

We primarily used library methods to tune hyperparameters. For random forests and decision trees, we used `GridSearchCV` to tune max depth, criterion (GINI vs entropy), number of estimators (for RF), and other splitting/leaf parameters for decision trees, which checked every combination of hyperparameters. This did not really affect results, however. For neural networks, we used `find_lr` to find our initial learning rate, and used a learning rate scheduler to tune learning rate over the course of the training. We also, over the course of selecting the model, tried different batch sizes, hidden layer sizes, number of hidden layers, activation functions, and loss functions, all of which we "manually tuned" by checking the accuracy results after training.

# CSM148 Spotiflies Code

13 December 2024

## 1 Exploration, Grouping, and Prediction of Genre in Spotify Songs

### 1.1 Team Spotiflies: Joanna, Aaron, Aubrey, Kennedy, Aster, Ethan

GitHub Link: <https://github.com/ketexon/csm148-spotiflies>

The data set that we chose for the project was the Spotify dataset, which is a dataset with information about the popularity, genre, and several other musical qualities of over 140K songs on the Spotify music streaming platform. The link to the original dataset can be found here: <https://huggingface.co/datasets/maharshipandya/spotify-tracks-dataset>.

The main feature we're interested in with this dataset is song genre, and how it relates to other features. We want to see if the musical qualities of a song are strong influences of what genre it's categorized. The dataset has 114 different genres, ranging from emo to classical to industrial and more. We are interested in answering questions regarding what, if anything, in songs may be a predictor of what genre it is, if certain genres are more similar to each other than others, and more.

### 1.2 First Looks and Cleaning Our Data

We started our project by investigating our data for any obvious cleaning we have to do before EDA, such as missing data or unreasonable values.

```
[50]: %pip install pandas numpy matplotlib seaborn scikit-learn mlxtend
```

```
Requirement already satisfied: pandas in ./venv/lib/python3.12/site-packages  
(2.2.3)
```

```
Requirement already satisfied: numpy in ./venv/lib/python3.12/site-packages  
(2.1.3)
```

```
Requirement already satisfied: matplotlib in ./venv/lib/python3.12/site-  
packages (3.9.3)
```

```
Requirement already satisfied: seaborn in ./venv/lib/python3.12/site-packages  
(0.13.2)
```

```
Requirement already satisfied: scikit-learn in ./venv/lib/python3.12/site-  
packages (1.5.2)
```

```
Requirement already satisfied: mlxtend in ./venv/lib/python3.12/site-packages  
(0.23.3)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in  
./venv/lib/python3.12/site-packages (from pandas) (2.9.0.post0)
```

```
Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.12/site-
```

packages (from pandas) (2024.2)  
 Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.12/site-packages (from pandas) (2024.2)  
 Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.12/site-packages (from matplotlib) (1.3.1)  
 Requirement already satisfied: cycycler>=0.10 in ./venv/lib/python3.12/site-packages (from matplotlib) (0.12.1)  
 Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.12/site-packages (from matplotlib) (4.55.2)  
 Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.12/site-packages (from matplotlib) (1.4.7)  
 Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.12/site-packages (from matplotlib) (24.2)  
 Requirement already satisfied: pillow>=8 in ./venv/lib/python3.12/site-packages (from matplotlib) (11.0.0)  
 Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.12/site-packages (from matplotlib) (3.2.0)  
 Requirement already satisfied: scipy>=1.6.0 in ./venv/lib/python3.12/site-packages (from scikit-learn) (1.14.1)  
 Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.12/site-packages (from scikit-learn) (1.4.2)  
 Requirement already satisfied: threadpoolctl>=3.1.0 in ./venv/lib/python3.12/site-packages (from scikit-learn) (3.5.0)  
 Requirement already satisfied: six>=1.5 in ./venv/lib/python3.12/site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)

[notice] A new release of pip is

available: 24.2 -> 24.3.1

[notice] To update, run:

`pip install --upgrade pip`

Note: you may need to restart the kernel to use updated packages.

```

[90]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, permutation_test_score,
↳ StratifiedShuffleSplit, GridSearchCV
from sklearn.inspection import permutation_importance
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.metrics import (
    mean_squared_error, mean_absolute_error, r2_score,
    roc_curve, auc,
    classification_report, accuracy_score, confusion_matrix,
    adjusted_rand_score, silhouette_score
)
  
```

```

from sklearn.mixture import GaussianMixture
from mlxtend.evaluate import bias_variance_decomp

from sklearn.preprocessing import PolynomialFeatures, StandardScaler,
    LabelEncoder, LabelBinarizer
from sklearn.pipeline import make_pipeline

from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

import torch
import torch.nn as nn
import torch.optim as optim
import lightning as pl
import lightning.pytorch.callbacks as pl_callbacks
from torch.utils.data import DataLoader, Dataset, TensorDataset
from typing import cast

import os
is_linux = os.name == 'posix'

```

```

[52]: # Reading in the data
spotify = pd.read_csv("dataset.csv")
spotify

```

```

[52]:      Unnamed: 0      track_id      artists \
0          0  5SuOikwiRyPMVoIQDJUGSV      Gen Hoshino
1          1  4qPNDBW1i3p13qLCtOKi3A      Ben Woodward
2          2  1iJBSr7s7jYXzM8EGcbK5b  Ingrid Michaelson;ZAYN
3          3  6lfxq3CG4xtTiEg7opyCyx      Kina Grannis
4          4  5vjLSffimiIP26QG5WcN2K      Chord Overstreet
...
113995      113995  2C3TZjDRiAzdyViavDJ217      Rainy Lullaby
113996      113996  1hIz5L4IB9hN3WRYPOCGPw      Rainy Lullaby
113997      113997  6x8ZfSoqDjuNa5SVP5QjvX      Cesária Evora
113998      113998  2e6sXL2bYv4bSz6VTdnfLs      Michael W. Smith
113999      113999  2hETkH7c0fqmz3LqZDHZf5      Cesária Evora

      album_name \
0          Comedy
1          Ghost (Acoustic)

```



2 To Begin Again  
3 Crazy Rich Asians (Original Motion Picture Sou...  
4 Hold On  
...  
113995 #mindfulness - Soft Rain for Mindful Meditatio...  
113996 #mindfulness - Soft Rain for Mindful Meditatio...  
113997 Best Of  
113998 Change Your World  
113999 Miss Perfumado

	track_name	popularity	duration_ms	explicit	\
0	Comedy	73	230666	False	
1	Ghost - Acoustic	55	149610	False	
2	To Begin Again	57	210826	False	
3	Can't Help Falling In Love	71	201933	False	
4	Hold On	82	198853	False	
...	...	...	...	...	
113995	Sleep My Little Boy	21	384999	False	
113996	Water Into Light	22	385000	False	
113997	Miss Perfumado	22	271466	False	
113998	Friends	41	283893	False	
113999	Barbincor	22	241826	False	

	danceability	energy	...	loudness	mode	speechiness	acousticness	\
0	0.676	0.4610	...	-6.746	0	0.1430	0.0322	
1	0.420	0.1660	...	-17.235	1	0.0763	0.9240	
2	0.438	0.3590	...	-9.734	1	0.0557	0.2100	
3	0.266	0.0596	...	-18.515	1	0.0363	0.9050	
4	0.618	0.4430	...	-9.681	1	0.0526	0.4690	
...	...	...	...	...	...	...	...	
113995	0.172	0.2350	...	-16.393	1	0.0422	0.6400	
113996	0.174	0.1170	...	-18.318	0	0.0401	0.9940	
113997	0.629	0.3290	...	-10.895	0	0.0420	0.8670	
113998	0.587	0.5060	...	-10.889	1	0.0297	0.3810	
113999	0.526	0.4870	...	-10.204	0	0.0725	0.6810	

	instrumentalness	liveness	valence	tempo	time_signature	\
0	0.000001	0.3580	0.7150	87.917	4	
1	0.000006	0.1010	0.2670	77.489	4	
2	0.000000	0.1170	0.1200	76.332	4	
3	0.000071	0.1320	0.1430	181.740	3	
4	0.000000	0.0829	0.1670	119.949	4	
...	...	...	...	...	...	
113995	0.928000	0.0863	0.0339	125.995	5	
113996	0.976000	0.1050	0.0350	85.239	4	
113997	0.000000	0.0839	0.7430	132.378	4	
113998	0.000000	0.2700	0.4130	135.960	4	

113999	0.000000	0.0893	0.7080	79.198	4
--------	----------	--------	--------	--------	---

```

        track_genre
0         acoustic
1         acoustic
2         acoustic
3         acoustic
4         acoustic
...
113995  world-music
113996  world-music
113997  world-music
113998  world-music
113999  world-music

```

[114000 rows x 21 columns]

We decided to remove the “Unnamed: 0” column, since we can uniquely identify songs from their track\_id, so the counter values aren’t very useful.

```
[53]: spotify_clean = spotify.drop(columns=["Unnamed: 0"])
```

Then we looked into NA values, and found one observation. We decided to drop this row since along with missing values, it also has other unexpected/unreasonable statistics, such as popularity being 0, and a duration of 0.

```
[54]: print(spotify_clean.isna().sum())
```

```

track_id      0
artists       1
album_name    1
track_name    1
popularity    0
duration_ms   0
explicit      0
danceability  0
energy        0
key           0
loudness      0
mode          0
speechiness   0
acousticness  0
instrumentalness 0
liveness      0
valence       0
tempo         0
time_signature 0
track_genre   0
dtype: int64

```

```
[55]: # Looking specifically at the row of missing values
spotify_clean[spotify_clean.isna().any(axis=1)]
```

```
[55]:
```

	track_id	artists	album_name	track_name	popularity	\
65900	1kR4gIb7nGxHPI3D2ifs59	NaN	NaN	NaN	0	

	duration_ms	explicit	danceability	energy	key	loudness	mode	\
65900	0	False	0.501	0.583	7	-9.46	0	

	speechiness	acousticness	instrumentalness	liveness	valence	\
65900	0.0605	0.69	0.00396	0.0747	0.734	

	tempo	time_signature	track_genre
65900	138.391	4	k-pop

```
[56]: # Dropping the row
spotify_clean = spotify_clean.dropna()
```

To investigate any unreasonable/odd values, we decided to visualize each numeric variable with boxplots.


```
[57]: # String and categorical columns in the dataset
# these will be dropped or one-hot-encoded depending on the model
string_columns = ['track_id', 'artists', 'album_name', 'track_name']
categorical_columns = ['key', 'mode', 'time_signature', 'track_genre']

# helper function to remove all categorical and string columns
def numeric(df: pd.DataFrame) -> pd.DataFrame:
    df_number = df.select_dtypes(include=[np.number])
    return df_number.loc[:, ~df_number.columns.isin(categorical_columns)]
```

```
[58]: # Creating boxplots
numeric_spotify = numeric(spotify_clean)
columns = numeric_spotify.columns

numeric_spotify[columns[:len(columns) // 2]].plot.box(figsize = (14, 6),
    ↳subplots=True)
numeric_spotify[columns[len(columns) // 2:]].plot.box(figsize = (14, 6),
    ↳subplots=True)
```

```
[58]: speechiness      Axes(0.125,0.11;0.110714x0.77)
acousticness      Axes(0.257857,0.11;0.110714x0.77)
instrumentalness  Axes(0.390714,0.11;0.110714x0.77)
liveness          Axes(0.523571,0.11;0.110714x0.77)
valence           Axes(0.656429,0.11;0.110714x0.77)
tempo             Axes(0.789286,0.11;0.110714x0.77)
dtype: object
```



Based on the boxplots, there seem to be quite a few numeric variables with a significant number of outliers. However, we decided to mainly only focus on the songs in which these outliers would be unreasonable based on domain knowledge, including duration being too low, the tempo being too slow, and the time signature being 0.

First, we addressed the songs with duration being too low (the song is too short). We decided that a threshold for a song to be “too short” as a song that is shorter than 30 seconds.

```
[59]: short_songs = spotify_clean[spotify_clean['duration_ms'] < 30000]

print(f"Number of tracks with duration less than 30 seconds:␣
      ↳{len(short_songs)}")
short_songs
```

Number of tracks with duration less than 30 seconds: 16

[59]:

	track_id \
11398	1egJZfc8JBT2blFQ4clPKe
16288	1T5QvLF9l04H030ZQbaX9p
16292	5viwzFJxwRE10EUR7G6hiD
16856	5YKCM3jbJ8lqUXUwfU7KwZ
39233	1T5QvLF9l04H030ZQbaX9p
39236	5viwzFJxwRE10EUR7G6hiD
59306	3qSaeaXmtOuzkqe7DKgoiM
59310	6hsyfegVY5yklJneM40mWi
59434	1AsX7B48DFJZplJEwmhGpl
59458	1sayezH8bWoxHMAQCccCTi
59609	1oVrTBrCsM2eTE1G50yxY9
59711	787rIUDmWWBIfZXwHUeXXQ
59775	1HVjSh7scH1PaPiLjy2LEu
59812	380gh3rsHba83kXx13gbKs
66925	ODVMaexfdXDz19zUv1zKej
101159	6bg9fEHIulR3DYNyWG0Jz

	artists \
11398	Benjamin Britten;Steven Isserlis
16288	Robert Schumann;Pavel Nersessian
16292	Robert Schumann;Pavel Nersessian
16856	Wolfgang Amadeus Mozart;Ingrid Haebler
39233	Robert Schumann;Pavel Nersessian
39236	Robert Schumann;Pavel Nersessian
59306	Leila Bela
59310	Leila Bela
59434	Alireza Mashayekhi;Ata Ebtekar;The Iranian Orc...
59458	Leila Bela
59609	Leila Bela
59711	Leila Bela
59775	Leila Bela;Leila's Opera Class
59812	Leila Bela
66925	Dora The Explorer
101159	Traditional;Cappella Musicale di Santa Maria i...

	album_name \
11398	October Classical Playlist
16288	Schumann, Poulenc & Others: Piano Works (Live ...
16292	Schumann, Poulenc & Others: Piano Works (Live ...
16856	Mozart: The Complete Piano Sonatas
39233	Schumann, Poulenc & Others: Piano Works (Live ...
39236	Schumann, Poulenc & Others: Piano Works (Live ...
59306	Angra Manyu
59310	Angra Manyu

59434	Ornamentationism
59458	Angra Manyu
59609	Angra Manyu
59711	Angra Manyu
59775	Angra Manyu
59812	Angra Manyu
66925	Dora The Explorer
101159	Giovannini: Messa a Quattro Breve Concertata

	track_name	popularity	\
11398	Cello Suite No. 3, Op. 87: IX. Passacaglia (Ex...	0	
16288	Carnaval, Op. 9: No. 20, Pause (Live in Japan,...	0	
16292	Carnaval, Op. 9: No. 13, Estrella (Live in Jap...	0	
16856	Andante in C Major, K. 1a	0	
39233	Carnaval, Op. 9: No. 20, Pause (Live in Japan,...	0	
39236	Carnaval, Op. 9: No. 13, Estrella (Live in Jap...	0	
59306	V-7	0	
59310	The Exorcism Begins...	0	
59434	Aural Blue	0	
59458	V-3	0	
59609	Shatter	0	
59711	Breath Ritual	0	
59775	Screams for a Finale! (feat. Leila's Opera Class)	0	
59812	V-4	0	
66925	Backpack, Backpack!	8	
101159	Pax Domini	0	

	duration_ms	explicit	danceability	energy	key	loudness	mode	\
11398	22266	False	0.335	0.0593	11	-26.365	0	
16288	17826	False	0.372	0.2780	8	-16.882	1	
16292	23506	False	0.379	0.2370	5	-18.265	1	
16856	17453	False	0.467	0.0301	2	-28.518	0	
39233	17826	False	0.372	0.2780	8	-16.882	1	
39236	23506	False	0.379	0.2370	5	-18.265	1	
59306	21120	False	0.229	0.0577	8	-27.960	0	
59310	8586	False	0.000	0.0400	8	-29.714	0	
59434	24666	False	0.187	0.9750	1	-8.223	1	
59458	28026	False	0.612	0.1370	8	-31.953	1	
59609	21240	False	0.424	0.8690	9	-8.168	0	
59711	28946	False	0.359	0.1430	4	-30.401	1	
59775	15800	False	0.251	0.5080	5	-10.564	0	
59812	13386	False	0.000	0.2240	11	-22.196	1	
66925	24266	False	0.903	0.3940	2	-8.018	1	
101159	24000	False	0.358	0.0335	1	-28.683	1	

	speechiness	acousticness	instrumentalness	liveness	valence	\
11398	0.0430	0.9920	0.8690	0.1160	0.1950	

16288	0.0370	0.9850	0.9210	0.1640	0.9120
16292	0.0470	0.9930	0.8870	0.1440	0.4770
16856	0.0428	0.9950	0.9000	0.1240	0.0000
39233	0.0370	0.9850	0.9210	0.1640	0.9120
39236	0.0470	0.9930	0.8870	0.1440	0.4770
59306	0.1960	0.6260	0.9310	0.1080	0.2530
59310	0.0000	0.9280	0.9560	0.1150	0.0000
59434	0.2360	0.0431	0.9800	0.3570	0.1360
59458	0.7920	0.8480	0.0000	0.0868	0.3930
59609	0.0728	0.7070	0.0893	0.1170	0.0000
59711	0.4390	0.3930	0.0379	0.3930	0.0366
59775	0.3160	0.9690	0.9990	0.9520	0.0000
59812	0.0000	0.9700	0.0000	0.9070	0.0000
66925	0.1040	0.0047	0.0000	0.0830	0.5430
101159	0.0435	0.9810	0.0000	0.2510	0.2980

	tempo	time_signature	track_genre
11398	77.266	5	british
16288	89.032	1	classical
16292	116.093	4	classical
16856	84.375	4	classical
39233	89.032	1	german
39236	116.093	4	german
59306	172.897	4	iranian
59310	0.000	0	iranian
59434	96.548	4	iranian
59458	100.765	4	iranian
59609	135.107	4	iranian
59711	119.207	5	iranian
59775	184.051	3	iranian
59812	0.000	0	iranian
66925	109.974	4	kids
101159	95.684	3	sleep

We found that there are only 16 songs that are less than 30 seconds. Since there were so few, we decided to look manually into these songs and investigate whether or not these songs seem to be real songs. Turns out, the majority of them are concert pauses (empty tracks) or sound effects and sound clips. Since we don't want these tracks to affect our investigation of genres, we decided to drop them.

```
[60]: spotify_clean = spotify_clean[spotify_clean['duration_ms'] >= 30000]
```

Next we looked at implausible 0 values with time signature and tempo. For context, time signature should be expressed in a fraction, while tempo is in BPM (so 0 BPM doesn't make sense).

```
[61]: print("Tempos == 0: ", spotify_clean["tempo"].eq(0).sum())
      print("Time signatures == 0: ", spotify_clean["time_signature"].eq(0).sum())
```

```
Tempos == 0: 155
```

Time signatures == 0: 161

```
[62]: tempo_timesig_0_songs = spotify_clean[(spotify_clean['tempo'] == 0) |  
      ↪(spotify_clean['time_signature'] == 0)]  
tempo_timesig_0_songs
```

```
[62]:          track_id \  
2926    0jdfbvSdaWvxfAlD20TtNc  
4131    59gg6zQhSKGVnkT3hWAY3l  
4379    4acmzQsAeMJJa5sGFSog7fu  
4664    1Kb2DqjHRvOcT5xeWtz3t5  
26910   7HSc2wpHlXKI18SCZK7zsP  
...  
101993   6H0kAiSAFB84jX7dgEDWd6  
112172   0jdfbvSdaWvxfAlD20TtNc  
113428   5EYzrykQ95u0mepteDi9KT  
113688   2EnZf7wbFv7ST4CJ3EvNzT  
113856   6XsYJ0dwT2hRzp0Qles78F
```

```
          artists \  
2926          Yaşlı Amca  
4131          Max Richter;Lang Lang  
4379    Dario Marianelli;Jack Liebeck;Benjamin Wallfisch  
4664          Sylvain Chauveau  
26910   Benny Martin  
...  
101993          Rain Sounds  
112172          Yaşlı Amca  
113428    El Ruido Blanco;Soñoliento Juan;Mantra para Do...  
113688          El Ruido Blanco  
113856          -
```

```
          album_name \  
2926          Akşamüstü  
4131          Voyager - Essential Max Richter  
4379    Jane Eyre - Original Motion Picture Soundtrack  
4664          Des Plumes Dans La Tête  
26910   Here Comes the Sun (Piano Instrumental)  
...  
101993          Rain  
112172          Akşamüstü  
113428          Aire Acondicionado de Ruido Blanco  
113688    Ruido Blanco para el bebé: sonidos relajantes ...  
113856          -
```

```
          track_name  popularity  duration_ms \  
2926          Sanki Yapamadım          44          213198
```



4131	The Departure	64	151506
4379	The End of Childhood (feat. Jack Liebeck)	55	73266
4664	Ferme Les Yeux	53	68493
26910	Here Comes the Sun (Piano Instrumental)	18	203705
...	...	...	...
101993	Rain: Natural Recording	32	84219
112172	Sanki Yapamadım	44	213198
113428	Aire de verano	27	128000
113688	Ruido Rosa Puro - Una Hora Versión	24	3601693
113856	( )	22	302185

	explicit	danceability	energy	key	loudness	mode	speechiness	\
2926	False	0.442	0.56700	8	-6.346	0	0.0516	
4131	False	0.000	0.03620	0	-22.519	0	0.0000	
4379	False	0.000	0.04450	0	-26.440	0	0.0000	
4664	False	0.000	0.03230	2	-23.636	0	0.0000	
26910	False	0.329	0.06070	9	-28.310	1	0.0507	
...	...	...	...	...	...	...	...	
101993	False	0.000	0.02540	8	-19.925	1	0.0000	
112172	False	0.442	0.56700	8	-6.346	0	0.0516	
113428	False	0.000	0.18800	8	-25.837	0	0.0000	
113688	False	0.000	0.00002	1	-11.165	1	0.0000	
113856	False	0.000	0.22400	8	-10.224	1	0.0000	

	acousticness	instrumentalness	liveness	valence	tempo	\
2926	0.238000	0.000325	0.0852	0.639	138.616	
4131	0.994000	0.940000	0.0958	0.000	0.000	
4379	0.972000	0.972000	0.0873	0.000	0.000	
4664	0.994000	0.973000	0.0922	0.000	0.000	
26910	0.994000	0.880000	0.0858	0.421	93.948	
...	...	...	...	...	...	
101993	0.000002	0.838000	0.3390	0.000	0.000	
112172	0.238000	0.000325	0.0852	0.639	138.616	
113428	0.139000	0.339000	0.1220	0.000	0.000	
113688	0.186000	1.000000	0.3620	0.000	0.000	
113856	0.142000	0.986000	0.4110	0.000	0.000	

	time_signature	track_genre
2926	0	alt-rock
4131	0	ambient
4379	0	ambient
4664	0	ambient
26910	0	disney
...	...	...
101993	0	sleep
112172	0	turkish
113428	0	world-music

```
113688          0  world-music
113856          0  world-music
```

```
[161 rows x 20 columns]
```

We manually looked through the songs with time signature and tempos of 0, and found that most of them tended to be ambient tracks, such as tracks of rainfall sounds or white noise. We decided that it might be good to remove these rows from our cleaned dataset, as we didn't want the zero values to skew the rest of our exploration in any way, and also because we hoped to look more at typical songs rather than ambient tracks.

Additionally, the overall count of these suspicious rows wasn't very high compared to size of the entire dataset (163 vs > 114k), so we figured that removing these values wouldn't drastically affect the size and comprehensiveness of our dataset.

We also confirmed this decision by seeing how the means of some of the other variables would change based on if we removed or kept these suspicious rows (see below).

```
[63]: spotify_clean_temp = spotify_clean[spotify_clean['time_signature']!= 0]
      spotify_clean_temp = spotify_clean[spotify_clean['tempo']!= 0]

      mean_temp = spotify_clean_temp.mean(numeric_only=True)
      mean_clean = spotify_clean.mean(numeric_only=True)

      # Calculate the difference
      mean_difference = mean_clean - mean_temp
      mean_difference
```

```
[63]: popularity          0.006038
      duration_ms        -73.741857
      explicit           -0.000116
      danceability        -0.000772
      energy              -0.000706
      key                 -0.000157
      loudness            -0.018772
      mode                0.000072
      speechiness         -0.000115
      acousticness        0.000240
      instrumentalness    0.000676
      liveness            0.000327
      valence             -0.000646
      tempo               -0.166333
      time_signature      -0.005316
      dtype: float64
```

From this, the effects on the means by removing the suspicious rows doesn't seem to be too large, with the largest being duration, which is measured in milliseconds, so would actually be less than a second difference. Thus, we dropped the rows where time signature or tempo were 0, which was a total of 161 rows.

```
[64]: spotify_clean = spotify_clean_temp
spotify_clean
```

```
[64]:
```

	track_id	artists \
0	5Su0ikwiRyPMVoIQDJUGSV	Gen Hoshino
1	4qPNDBW1i3p13qLCtOKi3A	Ben Woodward
2	1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN
3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis
4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet
...	...	...
113995	2C3TZjDRiAzdyViavDJ217	Rainy Lullaby
113996	1hIz5L4IB9hN3WRYP0CGPw	Rainy Lullaby
113997	6x8ZfSoqDjuNa5SVP5QjvX	Cesária Evora
113998	2e6sXL2bYv4bSz6VTdnfLs	Michael W. Smith
113999	2hETkH7c0fqmz3LqZDHZf5	Cesária Evora

	album_name \
0	Comedy
1	Ghost (Acoustic)
2	To Begin Again
3	Crazy Rich Asians (Original Motion Picture Sou...
4	Hold On
...	...
113995	#mindfulness - Soft Rain for Mindful Meditatio...
113996	#mindfulness - Soft Rain for Mindful Meditatio...
113997	Best Of
113998	Change Your World
113999	Miss Perfumado

	track_name	popularity	duration_ms	explicit \
0	Comedy	73	230666	False
1	Ghost - Acoustic	55	149610	False
2	To Begin Again	57	210826	False
3	Can't Help Falling In Love	71	201933	False
4	Hold On	82	198853	False
...	...	...	...	...
113995	Sleep My Little Boy	21	384999	False
113996	Water Into Light	22	385000	False
113997	Miss Perfumado	22	271466	False
113998	Friends	41	283893	False
113999	Barbincor	22	241826	False

	danceability	energy	key	loudness	mode	speechiness	acousticness \
0	0.676	0.4610	1	-6.746	0	0.1430	0.0322
1	0.420	0.1660	1	-17.235	1	0.0763	0.9240
2	0.438	0.3590	0	-9.734	1	0.0557	0.2100
3	0.266	0.0596	0	-18.515	1	0.0363	0.9050

4	0.618	0.4430	2	-9.681	1	0.0526	0.4690
...	...	...	...	...	...	...	...
113995	0.172	0.2350	5	-16.393	1	0.0422	0.6400
113996	0.174	0.1170	0	-18.318	0	0.0401	0.9940
113997	0.629	0.3290	0	-10.895	0	0.0420	0.8670
113998	0.587	0.5060	7	-10.889	1	0.0297	0.3810
113999	0.526	0.4870	1	-10.204	0	0.0725	0.6810

	instrumentalness	liveness	valence	tempo	time_signature	\
0	0.000001	0.3580	0.7150	87.917		4
1	0.000006	0.1010	0.2670	77.489		4
2	0.000000	0.1170	0.1200	76.332		4
3	0.000071	0.1320	0.1430	181.740		3
4	0.000000	0.0829	0.1670	119.949		4
...	...	...	...	...	...	...
113995	0.928000	0.0863	0.0339	125.995		5
113996	0.976000	0.1050	0.0350	85.239		4
113997	0.000000	0.0839	0.7430	132.378		4
113998	0.000000	0.2700	0.4130	135.960		4
113999	0.000000	0.0893	0.7080	79.198		4

	track_genre
0	acoustic
1	acoustic
2	acoustic
3	acoustic
4	acoustic
...	...
113995	world-music
113996	world-music
113997	world-music
113998	world-music
113999	world-music

[113828 rows x 20 columns]

Since there don't seem to be anymore values in our dataset that are unreasonable or missing, we exported our cleaned dataset to start exploring and visualizing our data.

```
[65]: # Exporting our cleaned dataset for analysis
spotify_clean.to_csv('spotify.csv', index=False)
```

### 1.3 Exploratory Data Analysis (EDA)

In order to find out a little more about the general structure and features of our dataset, we did some exploratory data analysis. We started by looking at the shape and structure of the data, and then the correlations between the numeric values, and then also did some visualizations of categorical features.

```
[66]: # Read in the cleaned data
spotify = pd.read_csv('spotify.csv')
spotify
```

```
[66]:
```

	track_id	artists \
0	5Su0ikwiRyPMVoIQDJUgSV	Gen Hoshino
1	4qPNDBW1i3p13qLct0Ki3A	Ben Woodward
2	1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN
3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis
4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet
...	...	...
113823	2C3TZjDRiAzdyViavDJ217	Rainy Lullaby
113824	1hIz5L4IB9hN3WRYP0CGPw	Rainy Lullaby
113825	6x8ZfSoqDjuNa5SVP5QjvX	Cesária Evora
113826	2e6sXL2bYv4bSz6VTdnfLs	Michael W. Smith
113827	2hETkH7c0fqmz3LqZDHZf5	Cesária Evora

	album_name \
0	Comedy
1	Ghost (Acoustic)
2	To Begin Again
3	Crazy Rich Asians (Original Motion Picture Sou...
4	Hold On
...	...
113823	#mindfulness - Soft Rain for Mindful Meditatio...
113824	#mindfulness - Soft Rain for Mindful Meditatio...
113825	Best Of
113826	Change Your World
113827	Miss Perfumado

	track_name	popularity	duration_ms	explicit \
0	Comedy	73	230666	False
1	Ghost - Acoustic	55	149610	False
2	To Begin Again	57	210826	False
3	Can't Help Falling In Love	71	201933	False
4	Hold On	82	198853	False
...	...	...	...	...
113823	Sleep My Little Boy	21	384999	False
113824	Water Into Light	22	385000	False
113825	Miss Perfumado	22	271466	False
113826	Friends	41	283893	False
113827	Barbincor	22	241826	False

	danceability	energy	key	loudness	mode	speechiness	acousticness \
0	0.676	0.4610	1	-6.746	0	0.1430	0.0322
1	0.420	0.1660	1	-17.235	1	0.0763	0.9240
2	0.438	0.3590	0	-9.734	1	0.0557	0.2100

3	0.266	0.0596	0	-18.515	1	0.0363	0.9050
4	0.618	0.4430	2	-9.681	1	0.0526	0.4690
...	...	...	...	...	...	...	...
113823	0.172	0.2350	5	-16.393	1	0.0422	0.6400
113824	0.174	0.1170	0	-18.318	0	0.0401	0.9940
113825	0.629	0.3290	0	-10.895	0	0.0420	0.8670
113826	0.587	0.5060	7	-10.889	1	0.0297	0.3810
113827	0.526	0.4870	1	-10.204	0	0.0725	0.6810

	instrumentalness	liveness	valence	tempo	time_signature	\
0	0.000001	0.3580	0.7150	87.917	4	
1	0.000006	0.1010	0.2670	77.489	4	
2	0.000000	0.1170	0.1200	76.332	4	
3	0.000071	0.1320	0.1430	181.740	3	
4	0.000000	0.0829	0.1670	119.949	4	
...	...	...	...	...	...	...
113823	0.928000	0.0863	0.0339	125.995	5	
113824	0.976000	0.1050	0.0350	85.239	4	
113825	0.000000	0.0839	0.7430	132.378	4	
113826	0.000000	0.2700	0.4130	135.960	4	
113827	0.000000	0.0893	0.7080	79.198	4	

	track_genre
0	acoustic
1	acoustic
2	acoustic
3	acoustic
4	acoustic
...	...
113823	world-music
113824	world-music
113825	world-music
113826	world-music
113827	world-music

[113828 rows x 20 columns]

```
[67]: spotify.describe()
```

	popularity	duration_ms	danceability	energy	\
count	113828.000000	1.138280e+05	113828.000000	113828.000000	
mean	33.237384	2.281340e+05	0.567603	0.642140	
std	22.314498	1.062882e+05	0.172371	0.250761	
min	0.000000	3.008000e+04	0.051300	0.000020	
25%	17.000000	1.742130e+05	0.456000	0.473000	
50%	35.000000	2.130000e+05	0.580000	0.685000	
75%	50.000000	2.615970e+05	0.695000	0.854000	

max	100.000000	5.237295e+06	0.985000	1.000000
-----	------------	--------------	----------	----------

	key	loudness	mode	speechiness \
count	113828.000000	113828.000000	113828.000000	113828.000000
mean	5.309186	-8.238433	0.637488	0.084758
std	3.559470	4.991371	0.480728	0.105735
min	0.000000	-46.591000	0.000000	0.022100
25%	2.000000	-10.000250	0.000000	0.035900
50%	5.000000	-6.997000	1.000000	0.049000
75%	8.000000	-5.000000	1.000000	0.084600
max	11.000000	4.532000	1.000000	0.965000

	acousticness	instrumentalness	liveness	valence \
count	113828.000000	113828.000000	113828.000000	113828.000000
mean	0.314602	0.155314	0.213220	0.474737
std	0.332302	0.308835	0.189925	0.258831
min	0.000000	0.000000	0.009250	0.000000
25%	0.016900	0.000000	0.098000	0.261000
50%	0.168000	0.000041	0.132000	0.464000
75%	0.597000	0.047600	0.273000	0.683000
max	0.996000	1.000000	1.000000	0.995000

	tempo	time_signature
count	113828.000000	113828.000000
mean	122.317259	3.909460
std	29.653245	0.407685
min	30.200000	0.000000
25%	99.432750	4.000000
50%	122.024000	4.000000
75%	140.078000	4.000000
max	243.372000	5.000000

```
[68]: # Select only numeric columns (remove categorical ones) using numeric()  
      ↪function written in data cleaning section  
      # Removed categorical columns: key, time signature, mode  
      numeric_spotify = numeric(spotify)
```

We started by looking at the distributions of all of the numeric variables with a histogram.

```
[69]: # Create subplots  
      num_columns = len(numeric_spotify.columns)  
      num_rows = (num_columns // 3) + (num_columns % 3 > 0)  
      fig, axes = plt.subplots(num_rows, 3, figsize=(15, num_rows * 4))  
      axes = axes.flatten()  
  
      # Create a histogram for each variable  
      for i, column in enumerate(numeric_spotify.columns):
```

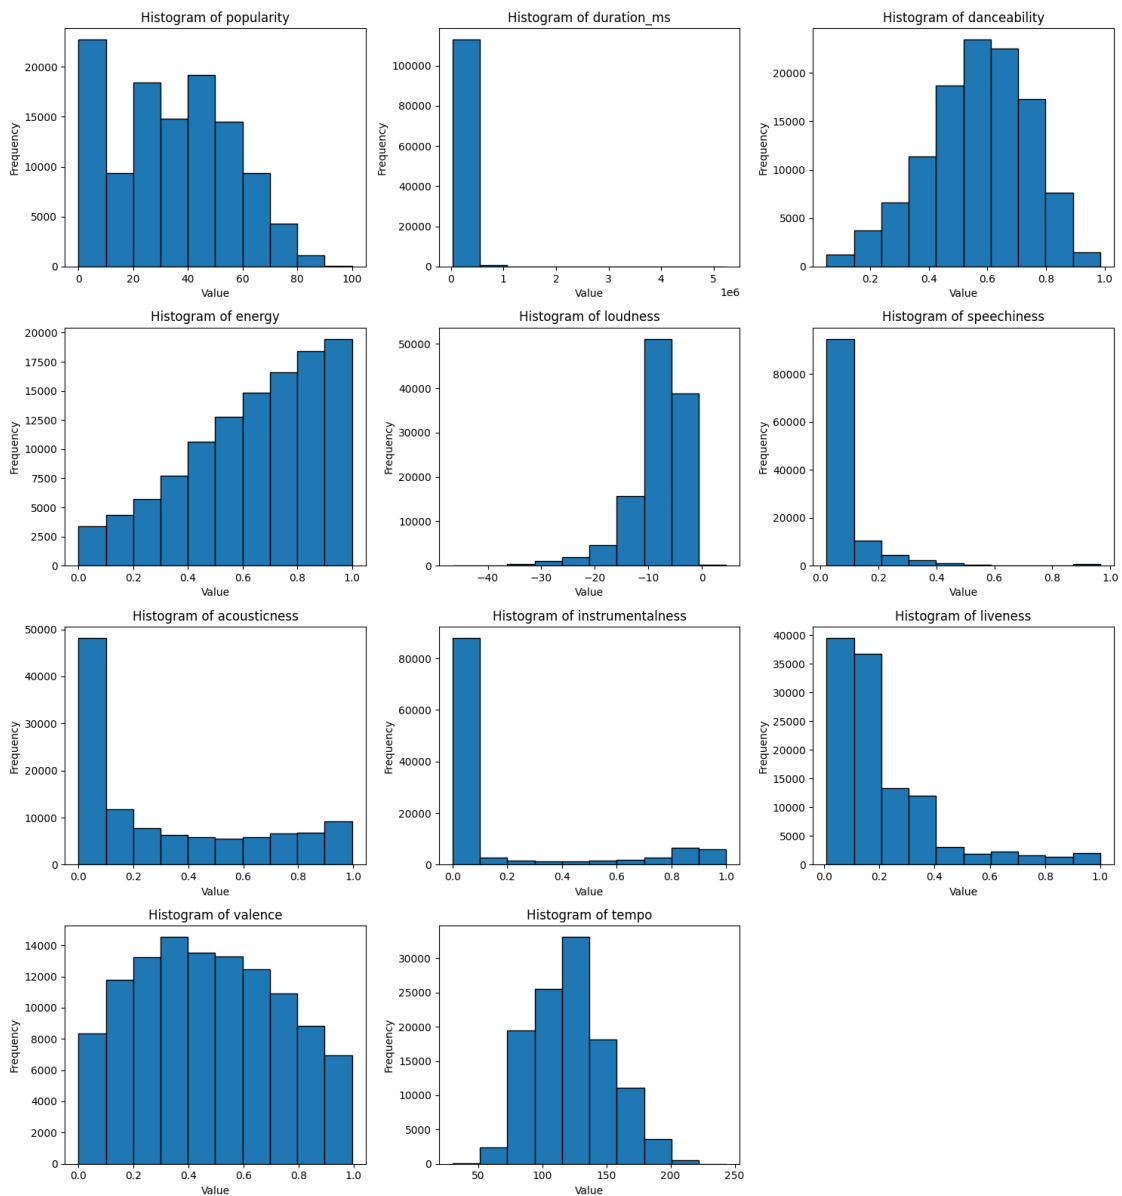
```

ax = axes[i]
numeric_spotify[column].hist(bins=10, edgecolor='black', ax=ax)
ax.set_title(f'Histogram of {column}')
ax.set_xlabel('Value')
ax.set_ylabel('Frequency')
ax.grid(False)

# Hiding unused subplots
for j in range(i + 1, len(axes)):
    axes[j].axis('off')

plt.tight_layout()
plt.show()

```





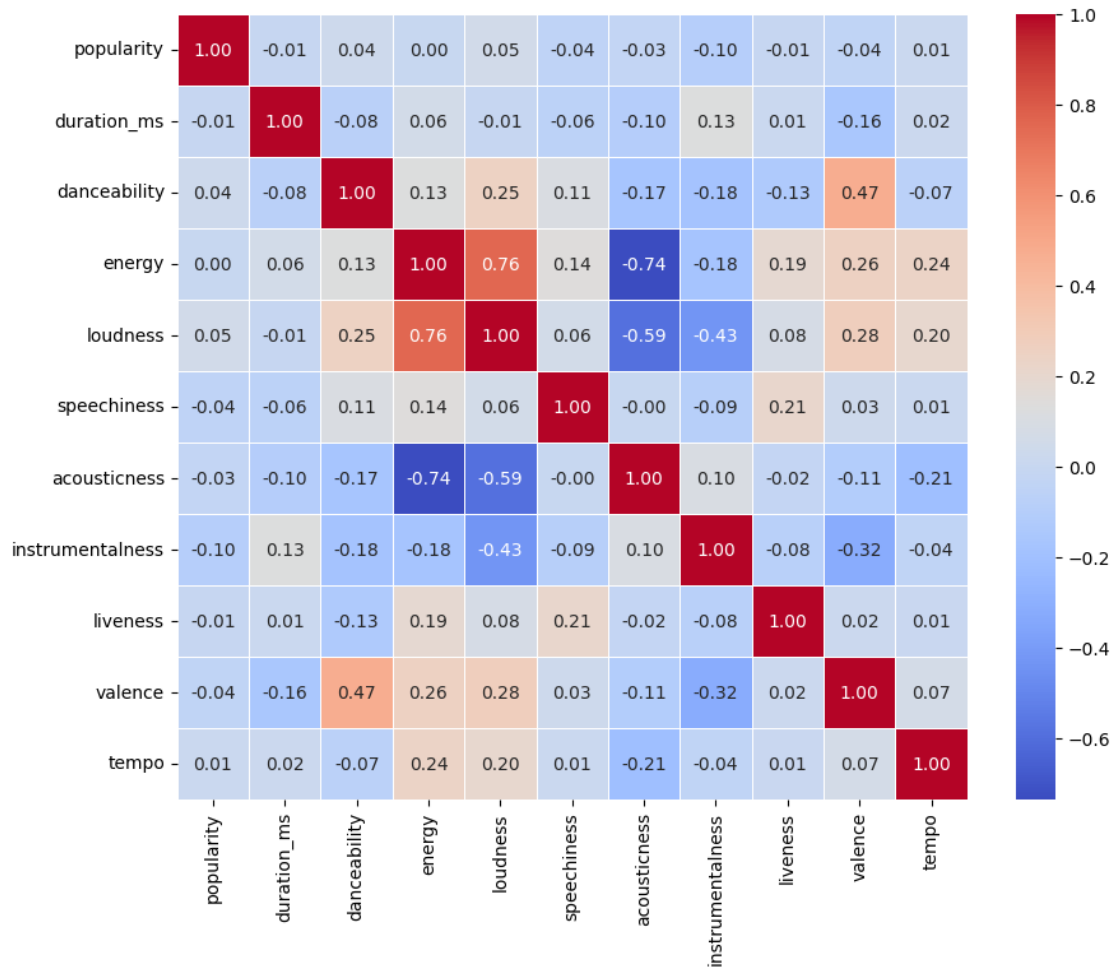
Based on this, we definitely have a few variables that are extremely skewed. However, we decided not to transform the data to make it more normal because the most of the data is already on a scale between 0 and 1, and we didn't want to lose possible meaningful information. For example, in the description of the data, it describes the “speechiness” variable with “Values above 0.66 describe tracks that are probably made entirely of spoken words. Values between 0.33 and 0.66 describe tracks that may contain both music and speech, either in sections or layered, including such cases as rap music. Values below 0.33 most likely represent music and other non-speech-like tracks”, which would be lost if we were to correct the right-skew of the speechiness variable.

We then decided to try to look at some relationships between variables by creating a correlation heatmap of every quantitative variable.

```
[70]: # Finding the correlation matrix
spotify_corr = numeric_spotify.corr()
spotify_corr

# Making correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(spotify_corr, annot=True, fmt='.2f', cmap='coolwarm', linewidths=0.
↪5)

plt.show()
```



Based on the heatmap of correlations, we wanted to see a more detailed scatterplots of the features that were more correlated. So, we decided to plot more detailed scatterplots of the correlations larger than 0.4 or smaller than -0.4. We then plotted a linear and quadratic regression for each.


```
[71]: # Filtering out the pairs of higher correlated features
correlated_features = [
    ("loudness", "energy"),
    ("acousticness", "energy"),
    ("acousticness", "loudness"),
    ("instrumentalness", "loudness"),
    ("valence", "danceability"),
]


# Making scatterplots for each variable vs popularity
for (i, j) in correlated_features:
```


```


plt.plot(np.unique(numeric_spotify[i]), np.poly1d(np.
↳polyfit(numeric_spotify[i], numeric_spotify[j], 1))(np.
↳unique(numeric_spotify[i])), color='red')
plt.plot(np.unique(numeric_spotify[i]), np.poly1d(np.
↳polyfit(numeric_spotify[i], numeric_spotify[j], 2))(np.
↳unique(numeric_spotify[i])), color='orange')
plt.scatter(numeric_spotify[i], numeric_spotify[j], s=5)
plt.xlabel(i)
plt.ylabel(j)
plt.show()


```











Though the variability from our fitted lines seemed pretty high on our scatterplots, some of the shapes of the graphs were pretty interesting. In particular, we thought the shapes of the (acousticness vs loudness), (instrumentalness vs loudness), (acousticness vs energy), and (energy vs loudness) graphs seemed to have a pretty distinct shape. The variables may have some sort of relationship with each other, which we noted in case they played some part in later sections of our exploration in our project.

We also looked into the categorical variables (excluding genre, which we will go into a more detailed exploration of later).

```
[72]: # Selecting the categories: key, mode, explicit, time signature
spotify_categories = spotify[categorical_columns + ['explicit']]
spotify_categories = spotify_categories.drop(columns='track_genre')

key_mapping = {
    0: 'C', 1: 'C /D ', 2: 'D', 3: 'D /E ', 4: 'E', 5: 'F', 6: 'F /G ', 7: 'G',
    8: 'G /A ', 9: 'A', 10: 'A /B ', 11: 'B', -1: 'No Key Detected'
}

# Map the 'key' column to pitch names
spotify_categories['key'] = spotify_categories['key'].map(key_mapping)
```

```

# Mode is 0 or 1 (1 for Major, 0 for Minor)
spotify_categories['mode'] = spotify_categories['mode'].map({0: 'Minor', 1: 'Major'})

# Time signature
spotify_categories['time_signature'] = spotify_categories['time_signature'].
    .astype(str)

spotify_categories

```

```

[72]:
      key  mode time_signature  explicit
0    C/D  Minor             4      False
1    C/D  Major             4      False
2      C  Major             4      False
3      C  Major             3      False
4      D  Major             4      False
...
113823  F  Major             5      False
113824  C  Minor             4      False
113825  C  Minor             4      False
113826  G  Major             4      False
113827  C/D  Minor             4      False

```

[113828 rows x 4 columns]

```

[73]: # Making barplots for each categorical value

# Create subplots for each categorical variable
plt.figure(figsize=(12, 10))
num_columns = len(spotify_categories.columns)
num_rows = (num_columns // 2) + (num_columns % 2 > 0)
fig, axes = plt.subplots(num_rows, 2, figsize=(12, num_rows * 5))
axes = axes.flatten()

# Loop through each categorical variable to create a bar plot
for i, column in enumerate(spotify_categories.columns):
    ax = axes[i]
    sns.countplot(x=column, data=spotify_categories, ax=ax)
    ax.set_title(f'Distribution of {column}')
    ax.set_xlabel(column)
    ax.set_ylabel('Count')
    ax.grid(False)

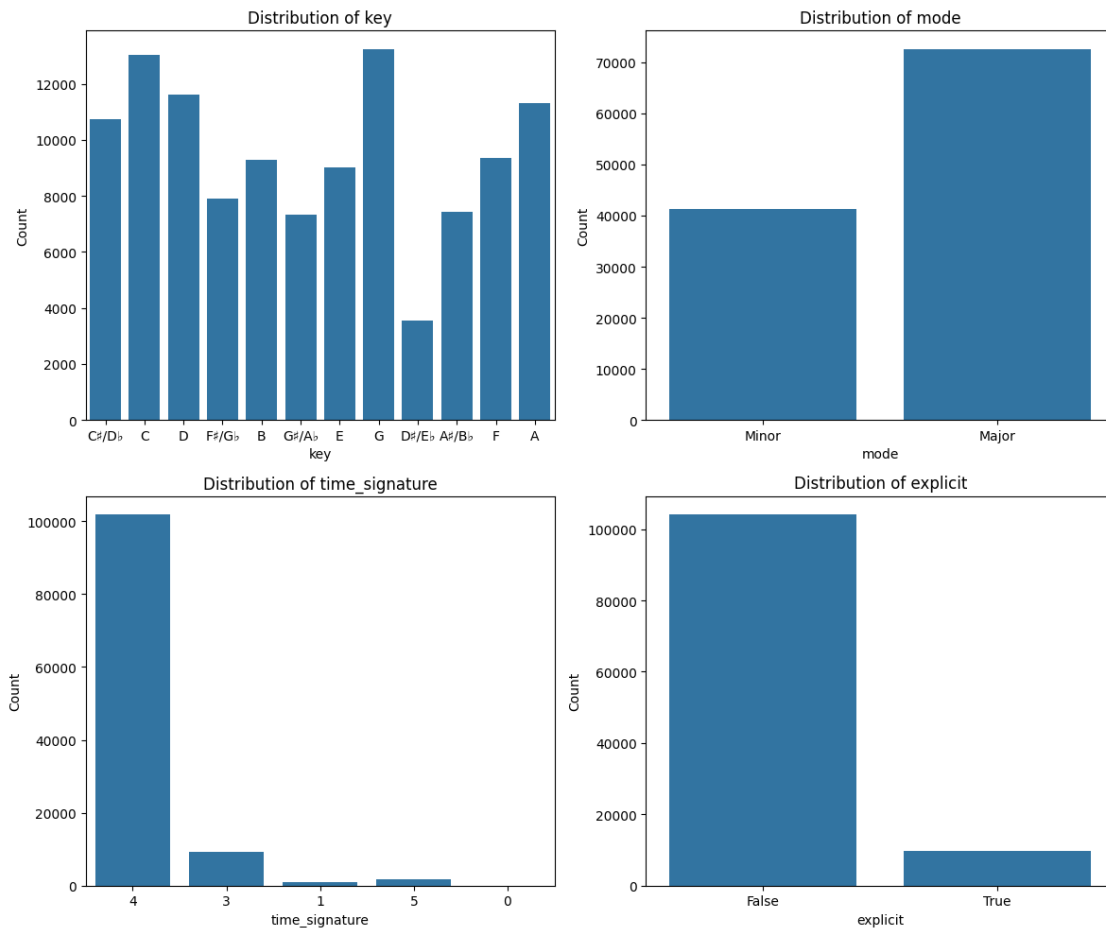
# Hiding unused subplots
for j in range(i + 1, len(axes)):
    axes[j].axis('off')

```



```
plt.tight_layout()
plt.show()
```

<Figure size 1200x1000 with 0 Axes>



Based on these visualizations it seems like the distributions for the key the song is in is fairly uniformly distributed, with D#/E the least number of tracks. Additionally, most songs seem to be in major, have a 4/4 time signature, or non-explicit.

After a more general exploration of some of the relationships and other features in our dataset, we decided to look closer into genre, the main focus of our project.

```
[74]: genre_counts = spotify['track_genre'].value_counts()
print(genre_counts)
```

```
track_genre
acoustic      1000
afrobeat      1000
alt-rock      1000
alternative    1000
```

```

anime          1000
...
classical      997
world-music    997
guitar         996
iranian        990
sleep          861
Name: count, Length: 114, dtype: int64

```

We have 114 different genres, and they all seem fairly split. In fact, our dataset seems to have 1000 samples for each genre, as the genres that have less than 1000 samples are genres that we cleaned data values from. We looked further into the specific genres and names.

```

[75]: genre_names = spotify['track_genre'].unique()
      print(genre_names)

['acoustic' 'afrobeat' 'alt-rock' 'alternative' 'ambient' 'anime'
 'black-metal' 'bluegrass' 'blues' 'brazil' 'breakbeat' 'british'
 'cantopop' 'chicago-house' 'children' 'chill' 'classical' 'club' 'comedy'
 'country' 'dance' 'dancehall' 'death-metal' 'deep-house' 'detroit-techno'
 'disco' 'disney' 'drum-and-bass' 'dub' 'dubstep' 'edm' 'electro'
 'electronic' 'emo' 'folk' 'forro' 'french' 'funk' 'garage' 'german'
 'gospel' 'goth' 'grindcore' 'groove' 'grunge' 'guitar' 'happy'
 'hard-rock' 'hardcore' 'hardstyle' 'heavy-metal' 'hip-hop' 'honky-tonk'
 'house' 'idm' 'indian' 'indie-pop' 'indie' 'industrial' 'iranian'
 'j-dance' 'j-idol' 'j-pop' 'j-rock' 'jazz' 'k-pop' 'kids' 'latin'
 'latino' 'malay' 'mandopop' 'metal' 'metalcore' 'minimal-techno' 'mpb'
 'new-age' 'opera' 'pagode' 'party' 'piano' 'pop-film' 'pop' 'power-pop'
 'progressive-house' 'psych-rock' 'punk-rock' 'punk' 'r-n-b' 'reggae'
 'reggaeton' 'rock-n-roll' 'rock' 'rockabilly' 'romance' 'sad' 'salsa'
 'samba' 'sertanejo' 'show-tunes' 'singer-songwriter' 'ska' 'sleep'
 'songwriter' 'soul' 'spanish' 'study' 'swedish' 'synth-pop' 'tango'
 'techno' 'trance' 'trip-hop' 'turkish' 'world-music']

```

We saw that some genres were a lot more specific than others, for example, the genre of “pop” versus “disney”, so we kept that in mind while doing later investigations and processing that we might want to group certain similar genres together, like grouping rock genres, pop genres, jazz, genres, etc.

We also looked at how some of the other variables were distributed grouped by track genre to see if certain genres had differentiable characteristics.

```

[76]: temp = spotify['track_genre']
      by_genres = numeric(spotify)
      by_genres['track_genre'] = temp

      by_genres.groupby('track_genre').mean()

```

```
[76]:
```

	popularity	duration_ms	danceability	energy	loudness	\
track_genre						
acoustic	42.483000	214896.957000	0.549593	0.435368	-9.447843	
afrobeat	24.399000	248412.791000	0.669580	0.702812	-7.789353	
alt-rock	33.943000	235455.907000	0.534493	0.754173	-6.191489	
alternative	24.337000	222016.180000	0.559927	0.720030	-6.078777	
ambient	44.151454	237478.207623	0.368974	0.237762	-18.584524	
...	...	...	...	...	...	
techno	39.042000	312311.477000	0.684348	0.746413	-8.077874	
trance	37.635000	269007.478000	0.583409	0.845272	-6.329711	
trip-hop	34.460000	274954.026000	0.634695	0.622363	-9.239915	
turkish	40.698000	219529.010000	0.616077	0.609804	-8.224722	
world-music	41.925777	294045.881645	0.415819	0.534177	-9.398517	

	speechiness	acousticness	instrumentalness	liveness	valence	\
track_genre						
acoustic	0.043247	0.566816	0.038336	0.153244	0.424023	
afrobeat	0.086579	0.270860	0.253483	0.184596	0.698619	
alt-rock	0.055071	0.122162	0.054097	0.210249	0.518260	
alternative	0.070101	0.147820	0.038159	0.201376	0.495570	
ambient	0.041687	0.776158	0.675362	0.129396	0.168002	
...	...	...	...	...	...	
techno	0.064212	0.081414	0.540038	0.159434	0.321878	
trance	0.079705	0.035870	0.423501	0.234357	0.276881	
trip-hop	0.076303	0.225615	0.383761	0.190342	0.478069	
turkish	0.105087	0.321125	0.035603	0.180750	0.462314	
world-music	0.041894	0.299980	0.089768	0.250248	0.251048	

	tempo
track_genre	
acoustic	119.010624
afrobeat	119.213337
alt-rock	124.634404
alternative	122.232394
ambient	111.447471
...	...
techno	128.255482
trance	133.276726
trip-hop	118.743616
turkish	120.367607
world-music	121.758988

```
[114 rows x 11 columns]
```

In order to better visualize these results, we decided to group the genres even further (as visualizing the differences between 114 grouped genres may be unclear). To do so, we manually grouped the genres into 10 more over-arching, general genres:

- **Pop:** cantopop, j-pop, k-pop, mandopop, pop, indie-pop, power-pop, pop-film, synth-pop
- **Rock:** alt-rock, alternative, hard-rock, punk-rock, psych-rock, rock, rock-n-roll, grunge, emo, rockabilly, guitar
- **Metal:** black-metal, death-metal, heavy-metal, metal, metalcore, grindcore
- **Electronic:** edm, electro, electronic, house, garage, techno, trance, dubstep, idm, minimal-techno, progressive-house, chicago-house, deep-house, detroit-techno, disco, drum-and-bass, dub, club, dance, dancehall
- **Hip-Hop:** hip-hop, rap, r-n-b, breakbeat
- **Jazz:** jazz, blues, soul, funk
- **Classical:** classical, opera, piano
- **World:** afrobeat, brazil, british, latin, latino, samba, salsa, reggae, reggaeton, tango, world-music, indian, iranian, turkish, malay, mpb, pagode, forro, french, german, spanish, swedish
- **Folk:** folk, bluegrass, country, singer-songwriter, songwriter, honky-tonk
- **Misc:** acoustic, ambient, anime, children, chill, comedy, disney, happy, party, study, sleep, show-tunes, new-age, kids, industrial, gospel, goth, groove, hardcore, hardstyle, indie, j-dance, j-idol, j-rock, punk, romance, sad, sertanejo, ska, trip-hop

```
[77]: original_data = pd.read_csv('spotify.csv')

genre_groups = {
    'pop': ['cantopop', 'j-pop', 'j-idol', 'k-pop', 'mandopop', 'pop',
    ↪ 'indie-pop', 'power-pop', 'pop-film', 'synth-pop'],
    'rock': ['alt-rock', 'alternative', 'hard-rock', 'indie', 'punk', 'j-rock',
    ↪ 'punk-rock', 'psych-rock', 'rock', 'rock-n-roll', 'grunge', 'emo',
    ↪ 'rockabilly', 'guitar'],
    'metal': ['black-metal', 'death-metal', 'heavy-metal', 'metal',
    ↪ 'metalcore', 'grindcore'],
    'electronic': ['edm', 'electro', 'electronic', 'house', 'garage',
    ↪ 'j-dance', 'hardcore', 'hardstyle', 'industrial', 'techno', 'trance',
    ↪ 'dubstep', 'idm', 'minimal-techno', 'progressive-house', 'chicago-house',
    ↪ 'deep-house', 'detroit-techno', 'disco', 'drum-and-bass', 'dub', 'club',
    ↪ 'dance', 'dancehall'],
    'hip-hop': ['hip-hop', 'rap', 'r-n-b', 'breakbeat', 'trip-hop'],
    'jazz': ['jazz', 'blues', 'soul', 'funk', 'ska', 'gospel'],
    'classical': ['classical', 'opera', 'piano'],
```

```

    'world': ['afrobeat', 'brazil', 'sertanejo', 'british', 'latin', 'latino',
↳ 'samba', 'salsa', 'reggae', 'reggaeton', 'tango', 'world-music', 'indian',
↳ 'iranian', 'turkish', 'malay', 'mpb', 'pagode', 'forro', 'french', 'german',
↳ 'spanish', 'swedish'],
    'folk': ['folk', 'bluegrass', 'country', 'singer-songwriter', 'songwriter',
↳ 'honky-tonk'],
    'misc': ['acoustic', 'ambient', 'anime', 'children', 'chill', 'comedy',
↳ 'disney', 'happy', 'party', 'study', 'sleep', 'show-tunes', 'new-age',
↳ 'kids', 'goth', 'groove', 'romance', 'sad']
}

genre_map = dict()
for genre, l in genre_groups.items():
    for original in l:
        genre_map[original] = genre

data = original_data.copy()
data['track_genre'] = data['track_genre'].map(genre_map)

data.to_csv(
    "grouped_cleaned_spotify.csv",
    index=False
)

```

```

[78]: grouped_spotify = pd.read_csv("grouped_cleaned_spotify.csv")
grouped_spotify['track_genre'].value_counts()

```

```

[78]: track_genre
electronic    24000
world         22984
misc          17855
rock          13996
pop           9999
metal         6000
folk          6000
jazz          5999
hip-hop       4000
classical     2995
Name: count, dtype: int64

```

Though now we have some imbalance between the different genres, it is much easier for us to visualize the values. Looking again at the means:

```

[79]: temp = grouped_spotify['track_genre']
by_grouped_genres = numeric(grouped_spotify)
by_grouped_genres['track_genre'] = temp

```

```
grouped_genre_means = by_grouped_genres.groupby('track_genre').mean()
grouped_genre_means
```

```
[79]:
```

	popularity	duration_ms	danceability	energy	loudness	\
track_genre						
classical	27.671452	223765.185309	0.383762	0.275876	-16.901773	
electronic	32.682375	244199.154417	0.628392	0.748102	-6.786162	
folk	28.781833	214844.595833	0.557295	0.484704	-9.592910	
hip-hop	32.257250	255499.212000	0.657858	0.699074	-7.198840	
jazz	28.944157	226728.059177	0.568563	0.578545	-8.072706	
metal	30.441833	238321.365167	0.375544	0.893222	-5.626057	
misc	32.046766	202753.826211	0.544729	0.519480	-11.052483	
pop	41.354135	236116.865087	0.576878	0.647109	-7.206672	
rock	36.582809	216122.028079	0.531166	0.673804	-7.515469	
world	33.082884	231900.404760	0.600875	0.635310	-7.881580	

	speechiness	acousticness	instrumentalness	liveness	valence	\
track_genre						
classical	0.048077	0.806379	0.432656	0.176999	0.303223	
electronic	0.092085	0.122595	0.249558	0.190737	0.432817	
folk	0.046643	0.528443	0.048824	0.178637	0.517899	
hip-hop	0.095396	0.205441	0.200120	0.216372	0.534306	
jazz	0.074014	0.385527	0.028089	0.215736	0.539197	
metal	0.098494	0.017465	0.241315	0.251464	0.296137	
misc	0.123619	0.488473	0.250852	0.230444	0.411343	
pop	0.061161	0.313732	0.028029	0.185500	0.511655	
rock	0.062589	0.258902	0.072948	0.204043	0.530961	
world	0.082780	0.372472	0.082835	0.243521	0.547976	

	tempo
track_genre	
classical	110.627554
electronic	126.276463
folk	120.840527
hip-hop	123.145587
jazz	119.278760
metal	126.118669
misc	117.992933
pop	124.357350
rock	124.747857
world	120.740060

```
[80]: ranges = grouped_genre_means.max() - grouped_genre_means.min()
ranges
```

```
[80]: popularity          13.682683
duration_ms          52745.385789
```

```

danceability      0.282314
energy            0.617346
loudness          11.275715
speechiness       0.076976
acousticness      0.788914
instrumentalness  0.404627
liveness          0.074465
valence           0.251839
tempo            15.648908
dtype: float64

```

```

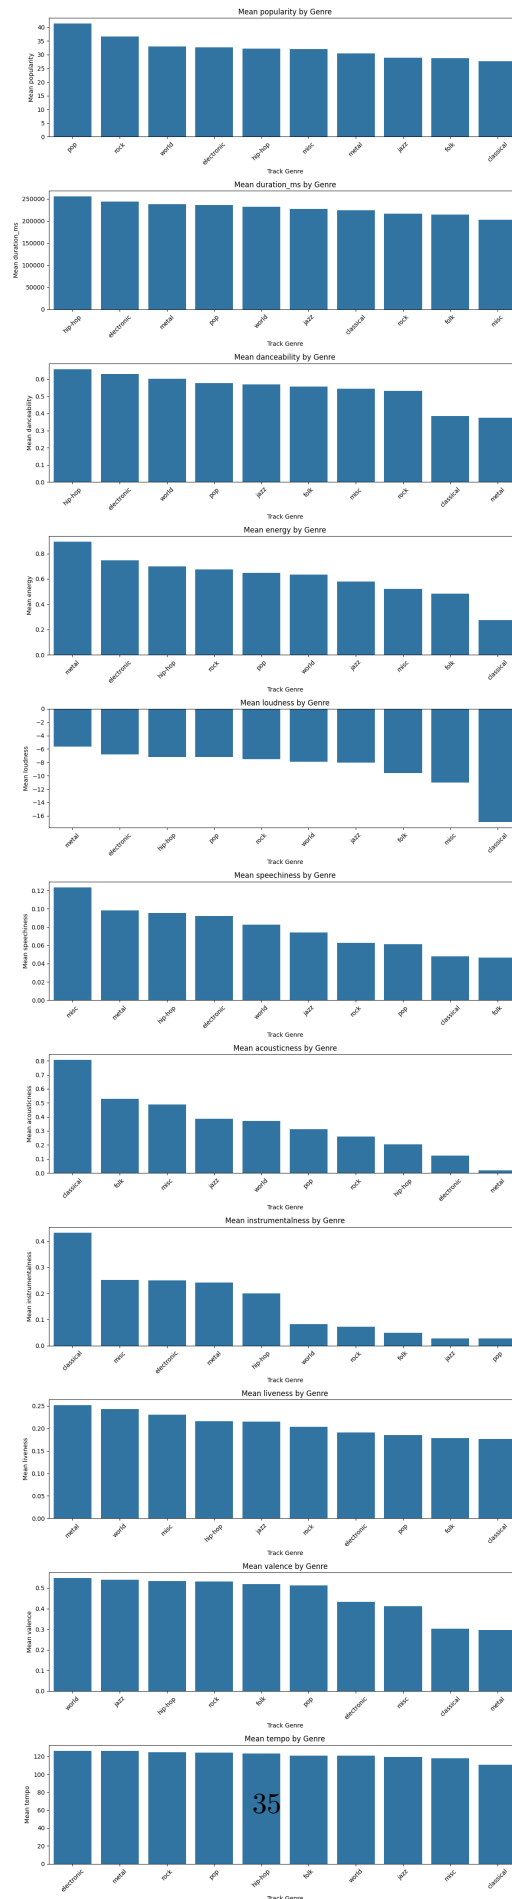
[81]: num_vars = len(grouped_genre_means.columns)
fig, axes = plt.subplots(nrows=num_vars, ncols=1, figsize=(12, num_vars * 4))

# Plot each variable as a barplot
for i, (column, ax) in enumerate(zip(grouped_genre_means.columns, axes)):
    sorted_data = grouped_genre_means[column].sort_values(ascending=False) # ↵
    ↪sort descending

    sns.barplot(
        x=sorted_data.index,
        y=sorted_data.values,
        ax=ax,
    )
    ax.set_title(f'Mean {column} by Genre')
    ax.set_xlabel('Track Genre')
    ax.set_ylabel(f'Mean {column}')
    ax.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

```





From the above visualizations, we can make a few interesting observations about each variable: - Pop music seems to be the most popular, though not necessarily by very much (around 5 points difference from rock). - Duration doesn't seem to have that much variability across genres when considering the range of the means is about 43000ms, which is about a 42 second difference. However, hip-hop songs seem to be the longest in duration compared to other genres. - Classical and metal tracks seem to score lower on danceability scores compared to other genres, and hip-hop is the highest for danceability. - Metal scores very high for energy levels, while classical seems to score lower. - Similarly, metal scores high for loudness, while classical scores much lower. In this variable, loudness, the difference between classical and the rest of the variables is much more apparent. - For the speechiness variable, folk and classical have the least amount of speechiness, while miscellaneous genres have the most. - For acousticness, classical music scores very high while metal scores extremely low. - With instrumentality, classical music once again scores high, with pop and jazz scoring low. - There doesn't seem to be too large of a difference in liveness across the variables, though metal does score the highest while classical scores the lowest. - Surprisingly, metal and classical both score low on valence, with jazz, hip-hop, and world tracks scoring higher. - Tempo seems to be fairly uniformly distributed across the genres, with classical and jazz barely scoring lower than the other genres.

In summary, it seems like classical tracks and metal tracks have the largest distinction between their variables, and the variables acousticness and instrumentality seem to have the most variance across genres.

To investigate our genres in our dataset further, such as which features may play a role in the labelling of genres and if there's a way for us to predict the genre of a song based on its features, we can prepare our data for further model and training. We decided we didn't need to do any feature engineering since several of our variables already ranged from 0 to 1, some with meaningful values (as discussed in the cleaning section). We also didn't add any extra data through data augmentation because we already had over 100k samples, which we felt would be sufficient in our investigation.

In order to prepare our data for more learning models, we split it into training, validation, and testing datasets (60-20-20). an example of how we split our data is in the code below, but we re-split it for each model for the sake of simplicity. We did this both for our original cleaned dataset and also our dataset grouped by our 10 generalized genres.

```
[82]: random_seed = 42

# encode categorical categories
for column in categorical_columns:
    spotify[column] = pd.Categorical(spotify[column])

# First split: separate out 20% for the test set
spotify_train_val, spotify_test = train_test_split(spotify, test_size=0.2,
    ↪random_state=random_seed)
# Second split: separate remaining 80% into 60% training and 20% validation
```

```

spotify_train, spotify_val = train_test_split(spotify_train_val, test_size=0.
↳25, random_state=random_seed) # 0.25 * 0.8 = 0.2

# Grouped genre dataset, grouped_spotify
# First split: separate out 20% for the test set
grouped_spotify_train_val, grouped_spotify_test =
↳train_test_split(grouped_spotify, test_size=0.2, random_state=random_seed)
# Second split: separate remaining 80% into 60% training and 20% validation
grouped_spotify_train, grouped_spotify_val =
↳train_test_split(grouped_spotify_train_val, test_size=0.25,
↳random_state=random_seed) # 0.25 * 0.8 = 0.2

```

A summary of the datasets that result from the train-validation-test splitting is below:

Original 114 genre dataset: - spotify\_train (60% of the total dataset) - spotify\_val (20% of the total dataset) - spotify\_test (20% of the total dataset)

Re-grouped 10 genre dataset: - grouped\_spotify\_train (60% of the total dataset) - grouped\_spotify\_val (20% of the total dataset) - grouped\_spotify\_test (20% of the total dataset)

```

[83]: # This code does OHE and removed genre from grouped_spotify
response_variable = "track_genre"
ohe_column_transformer_wo_genre = ColumnTransformer(
    transformers=[
        ('ohe', OneHotEncoder(sparse_output=False), [col for col in
↳categorical_columns if col != response_variable]),
        ('string', 'drop', string_columns),
        ('numeric', 'passthrough', [col for col in spotify.columns if col !=
↳response_variable and col not in string_columns])
    ],
    verbose_feature_names_out=False
)
ohe_column_transformer_wo_genre.set_output(transform="pandas")
ohe_column_transformer_wo_genre.fit(grouped_spotify)

ohe_grouped_column_transformer_wo_genre = ColumnTransformer(
    transformers=[
        ('ohe', OneHotEncoder(sparse_output=False), [col for col in
↳categorical_columns if col != response_variable]),
        ('string', 'drop', string_columns),
        ('numeric', 'passthrough', [col for col in grouped_spotify.columns if
↳col != response_variable and col not in string_columns])
    ],
    verbose_feature_names_out=False
)
ohe_grouped_column_transformer_wo_genre.set_output(transform="pandas")
ohe_grouped_column_transformer_wo_genre.fit(grouped_spotify)

```

```
[83]: ColumnTransformer(transformers=[('ohe', OneHotEncoder(sparse_output=False),
                                     ['key', 'mode', 'time_signature']),
                                     ('string', 'drop',
                                      ['track_id', 'artists', 'album_name',
                                       'track_name']),
                                     ('numeric', 'passthrough',
                                      ['popularity', 'duration_ms', 'explicit',
                                       'danceability', 'energy', 'key', 'loudness',
                                       'mode', 'speechiness', 'acousticness',
                                       'instrumentalness', 'liveness', 'valence',
                                       'tempo', 'time_signature'])],
      verbose_feature_names_out=False)
```

## 1.4 Linear + Polynomial Regression

For our linear regression, we chose to use the **energy** variable as our response variable, as it seemed to have the most relationships with other variables based on our correlation heatmap during our EDA, and we were interested in finding out what may contribute to the energy of a song.

To find out our predictor variables, we used forward feature selection to see which predictors may be most correlated with our response variable (energy).

To find out our predictor variables, we used forward feature selection to see which predictors may be most correlated with our response variable (energy).

```
[52]: response = "energy" # this was chosen based on correlation heatmap

# Making sure only numeric variables are used:
num_spotify_train = numeric(spotify_train)
num_spotify_val = numeric(spotify_val)
num_spotify_test = numeric(spotify_test)

# Reshape the data to fit the model
X_train = num_spotify_train.drop(columns=response)
y = num_spotify_train[response]

lin_reg = LinearRegression()

# Select features
selector = SequentialFeatureSelector(
    lin_reg,
    n_features_to_select='auto',
    direction='forward',
    scoring='r2',
    cv = 5
)

selector.fit(X_train, y)
```

```
selected_features = selector.get_feature_names_out(X_train.columns)
print("Selected Features: ", selected_features)
```

Selected Features: ['loudness' 'acousticness' 'instrumentalness' 'liveness' 'valence']

Based on this we found that 'loudness', 'acousticness', 'instrumentalness', 'liveness', and 'valence' were our selected features, and graphed these with energy as the response variable to investigate the relationships.

```
[53]: X = selector.transform(X_train)

X_test = selector.transform(num_spotify_test.drop(columns=response))
y_test = num_spotify_test[response]

X_val = selector.transform(num_spotify_val.drop(columns=response))
y_val = num_spotify_val[response]

lin_reg.fit(X, y)
y_pred = lin_reg.predict(X)


# Plotting scatter plot with a regression line
for feature in selected_features:
    # Scatter plot of the data points
    plt.scatter(X_train[feature], y, color='blue', s=5, label="Data Points")


    # Calculate the regression line
    feature_values = X_train[feature].values.reshape(-1, 1)
    temp_model = LinearRegression()
    temp_model.fit(feature_values, y)
    y_line = temp_model.predict(feature_values)


    # Plot the regression line
    plt.plot(feature_values, y_line, color='red', label="Regression Line")


    plt.xlabel(feature)
    plt.ylabel(response)
    plt.legend()

plt.show()
```














Based on these graphs, some of our correlations are likely not linear, and others likely have a large amount of variance. We calculated the evaluation metrics for our linear regression model to see how it was performing.

```
[54]: # Calculating Evaluation Metrics:
y_val_pred = lin_reg.predict(X_val) # Predict on validation set

# Calculate metrics for the training set
train_mse = mean_squared_error(y, y_pred) # Mean squared error
train_rmse = np.sqrt(train_mse) # Root mean squared error
train_mae = mean_absolute_error(y, y_pred) # Mean average error
train_mad = np.mean(np.abs(y - y_pred)) # Mean absolute deviation
train_r2 = r2_score(y, y_pred) # R2 (coefficient of determination)

# Calculate metrics for the validation set
val_mse = mean_squared_error(y_val, y_val_pred)
val_rmse = np.sqrt(val_mse)
val_mae = mean_absolute_error(y_val, y_val_pred)
val_mad = np.mean(np.abs(y_val - y_val_pred))
val_r2 = r2_score(y_val, y_val_pred)
```

```

# Print Results:
print(f"Training MSE: {train_mse}, rMSE: {train_rmse}, MAE: {train_mae}, MAD:␣
↳{train_mad}, R²: {train_r2}")
print(f"Validation MSE: {val_mse}, rMSE: {val_rmse}, MAE: {val_mae}, MAD:␣
↳{val_mad}, R²: {val_r2}")

# bias variance
avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
    lin_reg,
    X, y.values,
    X_test, y_test.values,
    loss='mse',
    random_seed=123
)

print(f"Loss, variance, and bias: {avg_expected_loss}, {avg_bias}, {avg_var}")

```

```

Training MSE: 0.01591975141780648, rMSE: 0.12617349728768906, MAE:
0.0986015100486982, MAD: 0.0986015100486982, R²: 0.7477655716969742
Validation MSE: 0.016297699800804963, rMSE: 0.1276624447549277, MAE:
0.09952336547066755, MAD: 0.09952336547066755, R²: 0.7385932898496084
Loss, variance, and bias: 0.01592652204991807, 0.015924634631379178,
1.8874185388953818e-06

```

Based on the evaluation metrics we calculated above, our model seems to fit the data decently. Most notably, the MSE is very low and close together for the training and validation sets, 0.016 for both. This suggests there is low error and the model fits the data.

The  $R^2$  values (0.743 and 0.746) are moderately strong, which indicates that the model is explaining a good amount of the variance in the dependent variable (energy). However, we know based on the graphs (that do not look very linear) that we could probably improve our model if we were to use a different model, such as a polynomial regression. However, since the exploration on energy doesn't necessarily contribute to our project goal, we decided to switch directions and look further into genre instead.

Though linear/polynomial regression isn't necessarily a model that aligns with our project goals (which focuses on investigating and predicting the genres, a categorical variable, or different songs), and thus we chose not to include it in our main report, we still did some light explorations of the relationships between our different variables for different genres using regression.

```

[55]: # Make a dictionary for the data for efficiency
genre_groups = {genre: group for genre, group in spotify.groupby('track_genre')}
grouped_genre_groups = {genre: group for genre, group in grouped_spotify.
↳groupby('track_genre')}

print(np.unique(spotify["track_genre"]))
print(np.unique(grouped_spotify['track_genre']))

```

```
['acoustic' 'afrobeat' 'alt-rock' 'alternative' 'ambient' 'anime'
```

```
'black-metal' 'bluegrass' 'blues' 'brazil' 'breakbeat' 'british'
'cantopop' 'chicago-house' 'children' 'chill' 'classical' 'club' 'comedy'
'country' 'dance' 'dancehall' 'death-metal' 'deep-house' 'detroit-techno'
'disco' 'disney' 'drum-and-bass' 'dub' 'dubstep' 'edm' 'electro'
'electronic' 'emo' 'folk' 'forro' 'french' 'funk' 'garage' 'german'
'gospel' 'goth' 'grindcore' 'groove' 'grunge' 'guitar' 'happy'
'hard-rock' 'hardcore' 'hardstyle' 'heavy-metal' 'hip-hop' 'honky-tonk'
'house' 'idm' 'indian' 'indie' 'indie-pop' 'industrial' 'iranian'
'j-dance' 'j-idol' 'j-pop' 'j-rock' 'jazz' 'k-pop' 'kids' 'latin'
'latino' 'malay' 'mandopop' 'metal' 'metalcore' 'minimal-techno' 'mpb'
'new-age' 'opera' 'pagode' 'party' 'piano' 'pop' 'pop-film' 'power-pop'
'progressive-house' 'psych-rock' 'punk' 'punk-rock' 'r-n-b' 'reggae'
'reggaeton' 'rock' 'rock-n-roll' 'rockabilly' 'romance' 'sad' 'salsa'
'samba' 'sertanejo' 'show-tunes' 'singer-songwriter' 'ska' 'sleep'
'songwriter' 'soul' 'spanish' 'study' 'swedish' 'synth-pop' 'tango'
'techno' 'trance' 'trip-hop' 'turkish' 'world-music']
['classical' 'electronic' 'folk' 'hip-hop' 'jazz' 'metal' 'misc' 'pop'
'rock' 'world']
```

```
[56]: # Selecting only the numeric features
numeric_spotify = spotify.select_dtypes(include=[np.number])
numeric_spotify = numeric_spotify.loc[:, ~numeric_spotify.columns.isin(["mode",
↪ "key", "time_signature"])]
numeric_spotify
```

```
[56]:
```

	popularity	duration_ms	danceability	energy	loudness	speechiness	\
0	73	230666	0.676	0.4610	-6.746	0.1430	
1	55	149610	0.420	0.1660	-17.235	0.0763	
2	57	210826	0.438	0.3590	-9.734	0.0557	
3	71	201933	0.266	0.0596	-18.515	0.0363	
4	82	198853	0.618	0.4430	-9.681	0.0526	
...	...	...	...	...	...	...	
113823	21	384999	0.172	0.2350	-16.393	0.0422	
113824	22	385000	0.174	0.1170	-18.318	0.0401	
113825	22	271466	0.629	0.3290	-10.895	0.0420	
113826	41	283893	0.587	0.5060	-10.889	0.0297	
113827	22	241826	0.526	0.4870	-10.204	0.0725	
...	...	...	...	...	...	...	
113823	0.0322	0.000001	0.3580	0.7150	87.917		
1	0.9240	0.000006	0.1010	0.2670	77.489		
2	0.2100	0.000000	0.1170	0.1200	76.332		
3	0.9050	0.000071	0.1320	0.1430	181.740		
4	0.4690	0.000000	0.0829	0.1670	119.949		
...	...	...	...	...	...		
113823	0.6400	0.928000	0.0863	0.0339	125.995		
113824	0.9940	0.976000	0.1050	0.0350	85.239		

113825	0.8670	0.000000	0.0839	0.7430	132.378
113826	0.3810	0.000000	0.2700	0.4130	135.960
113827	0.6810	0.000000	0.0893	0.7080	79.198

[113828 rows x 11 columns]

We'll start with using linear regression. We looked at the linear regression between each of the 114 genres and each other variable to see if there seemed to be any strong relationships. We investigated the  $R^2$ , variance, and bias for the fits.

```
[57]: numeric_dict = {
    'popularity': 0,
    'duration_ms': 0,
    'danceability': 0,
    'energy': 0,
    'loudness': 0,
    'speechiness': 0,
    'acousticness': 0,
    'instrumentalness': 0,
    'liveness': 0,
    'valence': 0,
    'tempo': 0
}
```

```
[ ]: # Modeling linear regression of genres to all other variables

highest_r2 = 0
highest_r2_genre = ''
highest_r2_response = ''

for genre in genre_groups:
    df = genre_groups.get(genre, pd.DataFrame())
    numeric_df = df.select_dtypes(include=[np.number])
    numeric_df = numeric_df.loc[:, ~numeric_df.columns.isin(["mode", "key", "time_signature"])]

    linear_reg = LinearRegression()

    # Select features
    selector = SequentialFeatureSelector(
        linear_reg,
        n_features_to_select='auto',
        direction='forward',
        scoring='r2',
        cv = 5
    )
```

```

max_train_r2 = 0
max_validation_r2 = 0
max_response = ''

for x in numeric_df.columns:
    random_seed = 42
    response = x
    # Splitting the data
    # First split: separate out 20% for the test set
    train_val, test = train_test_split(numeric_df, test_size=0.2,
    ↪random_state=random_seed)

    # Second split: separate remaining 80% into 60% training and 40%
    ↪validation
    train, val = train_test_split(train_val, test_size=0.25,
    ↪random_state=random_seed) # 0.25 * 0.8 = 0.2

    # Reshape the data to fit the model
    X_train = train.drop(columns=response)
    y = train[response]

    selector.fit(X_train, y)
    selected_features = selector.get_feature_names_out(X_train.columns)

    # Transform data sets
    X = selector.transform(X_train)

    X_test = selector.transform(test.drop(columns=response))
    y_test = test[response]

    X_val = selector.transform(val.drop(columns=response))
    y_val = val[response]

    linear_reg.fit(X, y)
    y_pred = linear_reg.predict(X)

    # Calculating Evaluation Metrics:
    y_val_pred = linear_reg.predict(X_val)

    # Calculate metrics for the training set
    train_mse = mean_squared_error(y, y_pred) # Mean squared error
    train_rmse = np.sqrt(train_mse) # Root mean squared error
    train_mae = mean_absolute_error(y, y_pred) # Mean average error
    train_mad = np.mean(np.abs(y - y_pred)) # Mean absolute deviation
    train_r2 = r2_score(y, y_pred) # R2 (coefficient of determination)

```

```

# Calculate metrics for the validation set
val_mse = mean_squared_error(y_val, y_val_pred)
val_rmse = np.sqrt(val_mse)
val_mae = mean_absolute_error(y_val, y_val_pred)
val_mad = np.mean(np.abs(y_val - y_val_pred))
val_r2 = r2_score(y_val, y_val_pred)

# Print Results:
print("")
print(f"{genre} x {response} - Training MSE: {train_mse}, rMSE:␣
↪{train_rmse}, MAE: {train_mae}, MAD: {train_mad}, R²: {train_r2}")
print(f"{genre} x {response} - Validation MSE: {val_mse}, rMSE:␣
↪{val_rmse}, MAE: {val_mae}, MAD: {val_mad}, R²: {val_r2}")

if val_r2 > max_validation_r2:
    max_train_r2 = train_r2
    max_validation_r2 = val_r2
    max_response = x

# bias variance
avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
    linear_reg,
    X, y.values,
    X_test, y_test.values,
    loss='mse',
    random_seed=123
)

print(f"{genre} x {response} - Loss, Variance, and Bias:␣
↪{avg_expected_loss}, {avg_bias}, {avg_var}")
print("")
print(f"Response Variable with the Highest Validation R² for {genre}:␣
↪{max_response}")
print(f"Train R² for {max_response} in {genre}: {max_train_r2}")
print(f"Validation R² for {max_response} in {genre}: {max_validation_r2}")
numeric_dict[max_response]+=1
print("")
print("-----")

# Storing the variable / genre with the highest R² values
if(max_validation_r2 > highest_r2):
    highest_r2 = max_validation_r2
    highest_r2_genre = genre
    highest_r2_response = response

print(numeric_dict)
print("")

```

```
print(f"The highest Validation R^2 was between {genre} and {response}, with a score of {highest_r2}.")
```

Output: The highest Validation R<sup>2</sup> was between world-music and tempo, with a score of 0.9132660013292169.

Based on our linear regression investigation, it seems like energy tended to have the highest validation correlation with various genres, with 101 genres that were correlated the most with energy compared to other features. Following energy, 11 genres were more correlated with loudness, and 1 genre for both speechiness and valence had highest correlations. The relationship that seemed to have the strongest linear correlation was between the world-music genre's variables, with the tempo variable as the response.

We decided to also try a quadratic regression on our data to see if the increase in degree may have any significant effects.

```
[59]: # Define the degree of the polynomial, in this case, quadratic
degree = 2

# Create a pipeline for Polynomial Regression
poly_reg = make_pipeline(
    PolynomialFeatures(degree=degree),
    LinearRegression()
)
```

```
[ ]: # Modeling quadratic regression (degree = 2) of genres to all other variables

numeric_dict2 = {
    'popularity' : 0,
    'duration_ms' : 0,
    'danceability' : 0,
    'energy' : 0,
    'loudness' : 0,
    'speechiness' : 0,
    'acousticness' : 0,
    'instrumentalness' : 0,
    'liveness' : 0,
    'valence' : 0,
    'tempo' : 0
}

highest_r2 = 0
highest_r2_genre = ''
highest_r2_response = ''

for genre in genre_groups:
    df = genre_groups.get(genre, pd.DataFrame())
    numeric_df = df.select_dtypes(include=[np.number])
```

```

numeric_df = numeric_df.loc[:, ~numeric_df.columns.isin(["mode", "key",
↪ "time_signature"])]

# Select features
selector = SequentialFeatureSelector(
    linear_reg,
    n_features_to_select='auto',
    direction='forward',
    scoring='r2',
    cv = 5
)

max_train_r2 = 0
max_validation_r2 = 0
max_response = ''

for x in numeric_df.columns:
    random_seed = 42
    response = x
    # Splitting the data
    # First split: separate out 20% for the test set
    train_val, test = train_test_split(numeric_df, test_size=0.2,
↪ random_state=random_seed)

    # Second split: separate remaining 80% into 60% training and 40%
↪ validation
    train, val = train_test_split(train_val, test_size=0.25,
↪ random_state=random_seed) # 0.25 * 0.8 = 0.2

    # Reshape the data to fit the model
    X_train = train.drop(columns=response)
    y = train[response]

    selector.fit(X_train, y)
    selected_features = selector.get_feature_names_out(X_train.columns)

    # Transform data sets
    X = selector.transform(X_train)

    X_test = selector.transform(test.drop(columns=response))
    y_test = test[response]

    X_val = selector.transform(val.drop(columns=response))
    y_val = val[response]

    # Fit the Polynomial Regression model

```



```

poly_reg.fit(X, y)
y_pred = poly_reg.predict(X)

# Calculating Evaluation Metrics:
y_val_pred = poly_reg.predict(X_val)

# Calculate metrics for the training set
train_mse = mean_squared_error(y, y_pred) # Mean squared error
train_rmse = np.sqrt(train_mse) # Root mean squared error
train_mae = mean_absolute_error(y, y_pred) # Mean average error
train_mad = np.mean(np.abs(y - y_pred)) # Mean absolute deviation
train_r2 = r2_score(y, y_pred) # R2 (coefficient of determination)

# Calculate metrics for the validation set
val_mse = mean_squared_error(y_val, y_val_pred)
val_rmse = np.sqrt(val_mse)
val_mae = mean_absolute_error(y_val, y_val_pred)
val_mad = np.mean(np.abs(y_val - y_val_pred))
val_r2 = r2_score(y_val, y_val_pred)

# Print Results:
print("")
print(f"{genre} x {response} - Training MSE: {train_mse}, rMSE:↵
↵{train_rmse}, MAE: {train_mae}, MAD: {train_mad}, R2: {train_r2}")
print(f"{genre} x {response} - Validation MSE: {val_mse}, rMSE:↵
↵{val_rmse}, MAE: {val_mae}, MAD: {val_mad}, R2: {val_r2}")

if val_r2 > max_validation_r2:
    max_train_r2 = train_r2
    max_validation_r2 = val_r2
    max_response = x

# bias variance
avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
    linear_reg,
    X, y.values,
    X_test, y_test.values,
    loss='mse',
    random_seed=123
)

print(f"{genre} x {response} - Loss, Variance, and Bias:↵
↵{avg_expected_loss}, {avg_bias}, {avg_var}")
print("")
print(f"Response Variable with the Highest Validation R2 for {genre}:↵
↵{max_response}")
print(f"Train R2 for {max_response} in {genre}: {max_train_r2}")

```

```

print(f"Validation R^2 for {max_response} in {genre}: {max_validation_r2}")
numeric_dict2[max_response]+=1
print("")
print("-----")

# Storing the variable / genre with the highest R^2 values
if(max_validation_r2 > highest_r2):
    highest_r2 = max_validation_r2
    highest_r2_genre = genre
    highest_r2_response = response

print(numeric_dict)
print("")
print(f"The highest Validation R^2 was between {genre} and {response}, with a
↪score of {highest_r2}.")

```

Output: The highest Validation R<sup>2</sup> was between world-music and tempo, with a score of 0.9293826016602357.

Compared to the linear regression earlier, the quadratic regression didn't seem to have much of a change at all, other than the highest validation R<sup>2</sup> value between the world-music genre's variables and tempo slightly increasing from 0.913 to 0.929.

We also tried linear regression with our 10 manually grouped genres rather than the original 114 to see if it would make a difference.

[ ]: *# Modeling linear regression of 10 generalized genres to all other variables*

```

numeric_dict3 = {
    'popularity' : 0,
    'duration_ms' : 0,
    'danceability' : 0,
    'energy' : 0,
    'loudness' : 0,
    'speechiness' : 0,
    'acousticness' : 0,
    'instrumentalness' : 0,
    'liveness' : 0,
    'valence' : 0,
    'tempo' : 0
}

highest_r2 = 0
highest_r2_genre = ''
highest_r2_response = ''

for genre in grouped_genre_groups:
    df = grouped_genre_groups.get(genre, pd.DataFrame())

```

```

numeric_df = df.select_dtypes(include=[np.number])
numeric_df = numeric_df.loc[:, ~numeric_df.columns.isin(["mode", "key",
↳ "time_signature"])]

linear_reg = LinearRegression()

# Select features
selector = SequentialFeatureSelector(
    linear_reg,
    n_features_to_select='auto',
    direction='forward',
    scoring='r2',
    cv = 5
)

max_train_r2 = 0
max_validation_r2 = 0
max_response = ''

for x in numeric_df.columns:
    random_seed = 42
    response = x
    # Splitting the data
    # First split: separate out 20% for the test set
    train_val, test = train_test_split(numeric_df, test_size=0.2,
↳ random_state=random_seed)

    # Second split: separate remaining 80% into 60% training and 40%
↳ validation
    train, val = train_test_split(train_val, test_size=0.25,
↳ random_state=random_seed) # 0.25 * 0.8 = 0.2

    # Reshape the data to fit the model
    X_train = train.drop(columns=response)
    y = train[response]

    selector.fit(X_train, y)
    selected_features = selector.get_feature_names_out(X_train.columns)

    # Transform data sets
    X = selector.transform(X_train)

    X_test = selector.transform(test.drop(columns=response))
    y_test = test[response]

    X_val = selector.transform(val.drop(columns=response))

```

```

y_val = val[response]

# Fit the Linear Regression model
linear_reg.fit(X, y)
y_pred = linear_reg.predict(X)

# Calculating Evaluation Metrics:
y_val_pred = linear_reg.predict(X_val)

# Calculate metrics for the training set
train_mse = mean_squared_error(y, y_pred) # Mean squared error
train_rmse = np.sqrt(train_mse) # Root mean squared error
train_mae = mean_absolute_error(y, y_pred) # Mean average error
train_mad = np.mean(np.abs(y - y_pred)) # Mean absolute deviation
train_r2 = r2_score(y, y_pred) #  $R^2$  (coefficient of determination)

# Calculate metrics for the validation set
val_mse = mean_squared_error(y_val, y_val_pred)
val_rmse = np.sqrt(val_mse)
val_mae = mean_absolute_error(y_val, y_val_pred)
val_mad = np.mean(np.abs(y_val - y_val_pred))
val_r2 = r2_score(y_val, y_val_pred)

# Print Results:
print("")
print(f"{genre} x {response} - Training MSE: {train_mse}, rMSE:↵
↵{train_rmse}, MAE: {train_mae}, MAD: {train_mad}, R²: {train_r2}")
print(f"{genre} x {response} - Validation MSE: {val_mse}, rMSE:↵
↵{val_rmse}, MAE: {val_mae}, MAD: {val_mad}, R²: {val_r2}")

if val_r2 > max_validation_r2:
    max_train_r2 = train_r2
    max_validation_r2 = val_r2
    max_response = x

# bias variance
avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
    linear_reg,
    X, y.values,
    X_test, y_test.values,
    loss='mse',
    random_seed=123
)

print(f"{genre} x {response} - Loss, Variance, and Bias:↵
↵{avg_expected_loss}, {avg_bias}, {avg_var}")
print("")

```

```

    print(f"Response Variable with the Highest Validation R^2 for {genre}: {max_response}")
    print(f"Train R^2 for {max_response} in {genre}: {max_train_r2}")
    print(f"Validation R^2 for {max_response} in {genre}: {max_validation_r2}")
    numeric_dict3[max_response] += 1
    print("")
    print("-----")

    # Storing the variable / genre with the highest R^2 values
    if(max_validation_r2 > highest_r2):
        highest_r2 = max_validation_r2
        highest_r2_genre = genre
        highest_r2_response = response

print(numeric_dict)
print("")
print(f"The highest Validation R^2 was between {genre} and {response}, with a score of {highest_r2}.")

```

Output: The highest Validation R<sup>2</sup> was between world and tempo, with a score of 0.8921830467528595.

Compared to the results of our linear regression with the 114 genres, most genres still seemed like energy as the response variable had the highest correlation once again. And similar to before, the highest correlation was within the world genre with tempo as the response variable, though the validation R<sup>2</sup> dropped in comparison to the linear regression with 114 genres (from 0.913 to 0.892).

Because we weren't necessarily trying to predict anything with our linear and polynomial regression models, and since most of them didn't seem to be doing so well that it suggested overfitting, we didn't think we needed to do regularization on our models. We used these models more for exploration rather than prediction.

However, if interested, we did do lasso and ridge regularization with our exploration of energy as a response variable, which can be found in the week 3 check-in.

## 1.5 Logistic Regression

For our logistic regression, we wanted to try to see if we could classify a track's genre based on other variables.

```

[111]: label_encoder = LabelEncoder()
label_encoder.fit(grouped_spotify[response_variable])
Y_train = grouped_spotify_train[response_variable]
Y_train = label_encoder.transform(Y_train)

Y_test = grouped_spotify_test[response_variable]
Y_test = label_encoder.transform(Y_test)

```

```

X_train = ohe_grouped_column_transformer_wo_genre.
    ↪transform(grouped_spotify_train)
X_test = ohe_grouped_column_transformer_wo_genre.transform(grouped_spotify_test)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# logistic regression
logreg = LogisticRegression(solver="lbfgs", max_iter=1000)
logreg.fit(X_train, Y_train)

```

[111]: LogisticRegression(max\_iter=1000)

```

[110]: # Evaluate the model
Y_pred = logreg.predict(X_test)
Y_pred_proba = logreg.predict_proba(X_test)

report = classification_report(
    Y_test,
    Y_pred,
    target_names=label_encoder.classes_,
)
print(report)

```

	precision	recall	f1-score	support
classical	0.36	0.11	0.16	607
electronic	0.44	0.68	0.54	4839
folk	0.31	0.18	0.23	1218
hip-hop	0.50	0.00	0.00	767
jazz	0.11	0.00	0.00	1242
metal	0.52	0.57	0.54	1203
misc	0.43	0.36	0.39	3504
pop	0.28	0.08	0.13	2020
rock	0.30	0.28	0.29	2828
world	0.34	0.52	0.41	4538
accuracy			0.39	22766
macro avg	0.36	0.28	0.27	22766
weighted avg	0.37	0.39	0.35	22766

Looking at the evaluation metrics for our logistic regression, it seems like the precision for our predictions is not very high, with the highest being jazz, at 52%.

```

[104]: # ROC and AUC
best_thresholds = np.zeros(Y_pred_proba.shape[1])

```

```

for i, genre in enumerate(label_encoder.classes_):
    fpr, tpr, thresholds = roc_curve(Y_test == i, Y_pred_proba[:, i])
    auc_score = auc(fpr, tpr)
    best_thresholds[i] = thresholds[np.argmax(tpr - fpr)]
    plt.plot(fpr, tpr, label=f"{genre} (AUC = {auc_score:.2f})")

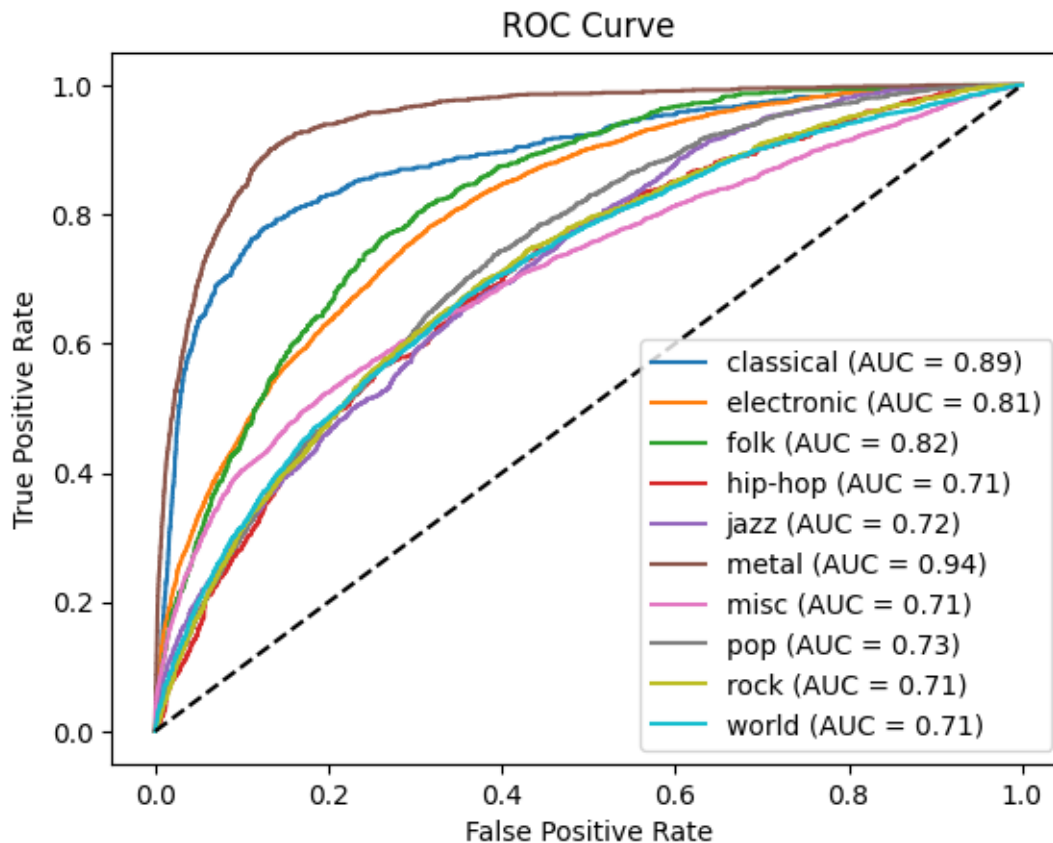
# print best threshold
best_threshold = best_thresholds.mean()
print(f"Best Threshold: {best_threshold}")

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()

```

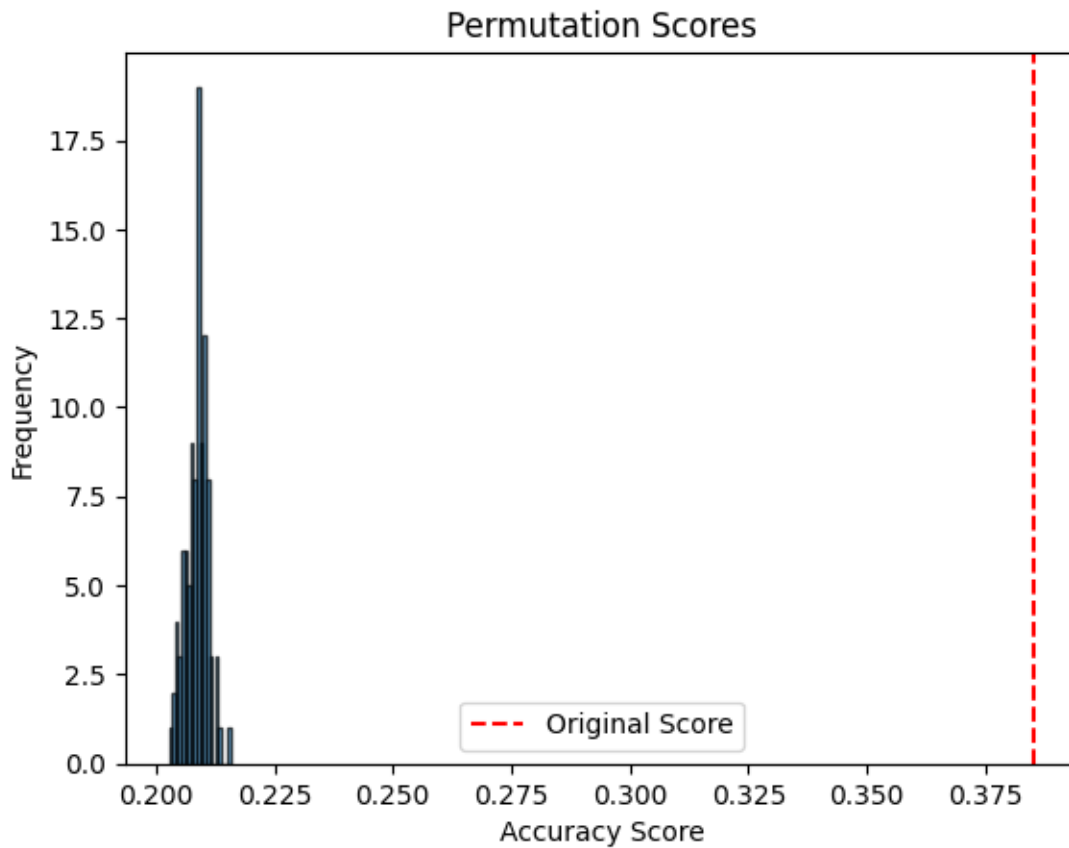
Best Threshold: 0.10861536576803565

[104]: <matplotlib.legend.Legend at 0x7d7c9c9d4710>



```
[87]: # permutation importance
perm_importance = permutation_test_score(
    logreg,
    X_test,
    Y_test,
    scoring="accuracy",
    cv=5,
    n_permutations=100,
    n_jobs=1,
    random_state=42
)
score, permutation_scores, p_value = perm_importance
```

```
[88]: # Plot permutation test scores
plt.hist(permutation_scores, bins=20, edgecolor='k', alpha=0.7)
plt.axvline(score, color='r', linestyle='--', label='Original Score')
plt.xlabel('Accuracy Score')
plt.ylabel('Frequency')
plt.title('Permutation Scores')
plt.legend()
plt.show()
```





```
[107]: # feature importances
feature_importances = permutation_importance(
    logreg,
    X_test,
    Y_test,
    n_repeats=50,
    n_jobs=1,
    random_state=42
)
pd.DataFrame(
    feature_importances.importances_mean,
    index=ohe_grouped_column_transformer_wo_genre.get_feature_names_out(),
    columns=["Importance"]
)
```

```
[107]:
```

	Importance
key_0	-0.000539
key_1	0.001236
key_2	0.000697
key_3	0.000407
key_4	0.000221
key_5	0.000479
key_6	-0.000257
key_7	0.001406
key_8	-0.000363
key_9	0.000019
key_10	0.000014
key_11	0.000649
mode_0	0.001288
mode_1	0.001288
time_signature_0	0.000000
time_signature_1	0.000172
time_signature_3	-0.000278
time_signature_4	-0.000046
time_signature_5	0.001314
popularity	0.003693
duration_ms	0.010849
explicit	0.002161
danceability	0.058037
energy	0.032968
key	-0.000184
loudness	0.005119
mode	0.001288
speechiness	0.013135
acousticness	0.066806

instrumentalness	0.020670
liveness	0.005648
valence	0.034678
tempo	0.003171
time_signature	-0.000537

## 1.6 KNN/Decision Trees/Random Forest

Since one of our main goals with the project was to see if we can classify or predict genres, KNN, decision trees, and random forests were good models for exploring this question. We decided to start by building a simple decision tree, using a grid search to tune the hyperparameters.

Decision trees are a good model for the data, since it is likely that a classification is not a simple function in any of the features, and cases overlap heavily.

```
[35]: response_variable = "track_genre"
Y_train = grouped_spotify_train[response_variable]
Y_test = grouped_spotify_test[response_variable]

X_train = grouped_spotify_train.drop(columns=[response_variable,
↳*string_columns])
X_test = grouped_spotify_test.drop(columns=[response_variable, *string_columns])
```

```
[ ]: # tuning decision tree hyperparameters
dt_parameters = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': [None, 'sqrt', 'log2']
}

dt_classifier = DecisionTreeClassifier(
    max_depth=50,
    random_state=42
)

grid_search = GridSearchCV(
    dt_classifier,
    param_grid=dt_parameters,
    cv=5,
    n_jobs=3,
    verbose=2
)

grid_search.fit(X_train, Y_train)
```

```
[ ]: GridSearchCV(cv=5,
                  estimator=DecisionTreeClassifier(max_depth=50, random_state=42),
                  n_jobs=3,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': [None, 10, 20, 30, 40, 50],
                              'max_features': [None, 'sqrt', 'log2'],
                              'min_samples_leaf': [1, 2, 4],
                              'min_samples_split': [2, 5, 10]})
```

```
[64]: dt_classifier = grid_search.best_estimator_

# predictions for testing
Y_pred = dt_classifier.predict(X_test)

# Evaluation metrics
dt_class_report = classification_report(Y_test, Y_pred, zero_division=1)
dt_conf_matrix = confusion_matrix(Y_test, Y_pred)

print(f'Classification Report: \n {dt_class_report}')
print(f'Confusion Matrix: \n {dt_conf_matrix}')
```

Classification Report:

	precision	recall	f1-score	support
classical	0.58	0.34	0.43	607
electronic	0.49	0.62	0.55	4026
folk	0.37	0.27	0.31	1218
hip-hop	0.35	0.09	0.14	578
jazz	0.38	0.27	0.32	830
metal	0.63	0.54	0.58	1203
misc	0.48	0.53	0.51	5931
pop	0.38	0.25	0.30	1825
rock	0.38	0.25	0.30	2227
world	0.44	0.56	0.49	4321
accuracy			0.46	22766
macro avg	0.45	0.37	0.39	22766
weighted avg	0.46	0.46	0.45	22766

Confusion Matrix:

```
[[ 206  26  13   0  18   0  208  29  24  83]
 [   2 2506  71  46  32  69  655 114 109 422]
 [   2  114 327  10  71   0  267  71  76 280]
 [   0  156   8  52  10  11  105  54  31 151]
 [   2   87  86   2 227   1  101  38  76 210]
 [   0   93   6   5   0 651  283   9 133  23]
 [  89  803 116  17  88 171 3172 170 263 1042]
 [   7  337  74   4  51  37  400 458  79 378]
```

```
[ 6 355 113 6 39 78 582 70 554 424]
[ 39 595 80 5 61 22 821 198 101 2399]]
```

Our decision tree seems to have about 46% accuracy, meaning it classified about 46% of the data correctly. This is better than chance, but could still be improved. In particular, it seems our decision tree was best at calculating electronic tracks, with a precision of 62%, while it was very bad at classifying hip-hop tracks, with a precision score of only 9%.

Based on the confusion matrix, it seems like electronic tracks, when misclassified, tended to be mislabeled as pop (276 instances) or rock (265 instances), which may be some indicator towards the similarities or overlaps with these genres. We also noticed that hip-hop tends to be mislabelled as electronic, as with hip-hop tracks only 160 were correctly labelled while 283 were labelled as electronic.

However, it is important to note that when we combined the genres into 10 more generalized genres, we caused an imbalance in samples. For example, electronic tracks have several more samples than hip-hop ones (20000 vs 3000) which may influence our decision tree's balance.

To learn more about which features in particular were contributing to our classification, we looked at the feature importances of our decision tree.

```
[66]: # Get feature importances
feature_importances = dt_classifier.feature_importances_
for i, importance in enumerate(feature_importances):
    print(f'{X_train.columns[i]}: {importance}')
```

```
popularity: 0.2200472160153845
duration_ms: 0.0852801213153109
explicit: 0.006176993344808496
danceability: 0.11891001330395011
energy: 0.04639268264190049
key: 0.0010495388111708972
loudness: 0.04128858458997459
mode: 0.006500681736710005
speechiness: 0.06524558872054981
acousticness: 0.18262725408120783
instrumentalness: 0.09890842233853107
liveness: 0.014775753707866195
valence: 0.05226490886822944
tempo: 0.05978402867389046
time_signature: 0.000748211850515286
```

```
[68]: # permutation importance
perm_importance = permutation_importance(
    dt_classifier,
    X_test,
    Y_test,
    n_repeats=20,
    random_state=42
)
```

```
perm_importance_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance Mean': perm_importance.importances_mean,
    'Importance Std': perm_importance.importances_std
})

perm_importance_df
```

```
[68]:
```

	Feature	Importance Mean	Importance Std
0	popularity	0.153321	0.002175
1	duration_ms	0.038263	0.001421
2	explicit	0.002715	0.000277
3	danceability	0.070403	0.001225
4	energy	0.026368	0.001298
5	key	0.000020	0.000083
6	loudness	0.021769	0.001009
7	mode	0.002462	0.000285
8	speechiness	0.024363	0.000918
9	acousticness	0.119643	0.002600
10	instrumentalness	0.058062	0.001181
11	liveness	0.004871	0.000420
12	valence	0.028775	0.001081
13	tempo	0.024633	0.000823
14	time_signature	0.000198	0.000110

Out of this, it seems like acousticness, popularity, and danceability have the highest importance when it comes to our decision tree. On the other hand, some of the time signatures, keys, and modes didn't really seem to be as key in feature importance.

To add more complexity to our model in hopes of improving its accuracy, we decided to try to train a random forest model on our data as well. We used a grid search again, but this time, we reduced the number of features in the search due to the higher complexity of the model requiring much longer time to do tuning.

```
[38]: rf_parameters = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20, 30],
}

# Random forest
rf_classifier = RandomForestClassifier(
    n_estimators=100,
    max_depth=20,
    random_state=42,
)

grid_search = GridSearchCV(
```

```

    estimator=rf_classifier,
    param_grid=rf_parameters,
    cv=5,
    n_jobs=1,
    verbose=2
)
grid_search.fit(X_train, Y_train)

```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```

[CV] END ...max_depth=10, n_estimators=100; total time= 6.3s
[CV] END ...max_depth=10, n_estimators=100; total time= 6.3s
[CV] END ...max_depth=10, n_estimators=100; total time= 6.3s
[CV] END ...max_depth=10, n_estimators=100; total time= 6.3s
[CV] END ...max_depth=10, n_estimators=100; total time= 6.2s
[CV] END ...max_depth=10, n_estimators=200; total time= 12.4s
[CV] END ...max_depth=10, n_estimators=200; total time= 12.6s
[CV] END ...max_depth=10, n_estimators=200; total time= 12.5s
[CV] END ...max_depth=10, n_estimators=200; total time= 12.5s
[CV] END ...max_depth=10, n_estimators=200; total time= 12.4s
[CV] END ...max_depth=20, n_estimators=100; total time= 10.8s
[CV] END ...max_depth=20, n_estimators=100; total time= 10.8s
[CV] END ...max_depth=20, n_estimators=100; total time= 10.9s
[CV] END ...max_depth=20, n_estimators=100; total time= 10.8s
[CV] END ...max_depth=20, n_estimators=100; total time= 10.9s
[CV] END ...max_depth=20, n_estimators=200; total time= 21.8s
[CV] END ...max_depth=20, n_estimators=200; total time= 21.1s
[CV] END ...max_depth=20, n_estimators=200; total time= 21.0s
[CV] END ...max_depth=20, n_estimators=200; total time= 20.9s
[CV] END ...max_depth=20, n_estimators=200; total time= 20.7s
[CV] END ...max_depth=30, n_estimators=100; total time= 11.0s
[CV] END ...max_depth=30, n_estimators=100; total time= 11.1s
[CV] END ...max_depth=30, n_estimators=100; total time= 11.4s
[CV] END ...max_depth=30, n_estimators=100; total time= 11.5s
[CV] END ...max_depth=30, n_estimators=100; total time= 11.2s
[CV] END ...max_depth=30, n_estimators=200; total time= 22.3s
[CV] END ...max_depth=30, n_estimators=200; total time= 21.9s
[CV] END ...max_depth=30, n_estimators=200; total time= 22.0s
[CV] END ...max_depth=30, n_estimators=200; total time= 22.3s
[CV] END ...max_depth=30, n_estimators=200; total time= 22.5s

```

```

[38]: GridSearchCV(cv=5,
    estimator=RandomForestClassifier(max_depth=20, random_state=42),
    n_jobs=1,
    param_grid={'max_depth': [10, 20, 30], 'n_estimators': [100, 200]},
    verbose=2)

```

```
[41]: print(grid_search.best_params_)
      rf_classifier = grid_search.best_estimator_
      Y_pred_rf = rf_classifier.predict(X_test)
```

```
{'max_depth': 20, 'n_estimators': 200}
```

```
[46]: # Evaluation metrics
      rf_train_accuracy = accuracy_score(Y_train, rf_classifier.predict(X_train))
      rf_accuracy = accuracy_score(Y_test, Y_pred_rf)
      rf_class_report = classification_report(Y_test, Y_pred_rf, zero_division=1)
      rf_conf_matrix = confusion_matrix(Y_test, Y_pred_rf)

      print(f'Train accuracy: {rf_train_accuracy}')
      print(f'Accuracy: {rf_accuracy}')
      print(f'Classification Report: \n {rf_class_report}')
      print(f'Confusion Matrix:')
      rf_conf_matrix
```

Train accuracy: 0.9188239428370621

Accuracy: 0.5578054994289731

Classification Report:

	precision	recall	f1-score	support
classical	0.60	0.36	0.45	607
electronic	0.66	0.70	0.68	4026
folk	0.68	0.55	0.61	1218
hip-hop	0.30	0.08	0.12	578
jazz	0.60	0.35	0.45	830
metal	0.65	0.62	0.63	1203
misc	0.50	0.65	0.57	5931
pop	0.48	0.34	0.40	1825
rock	0.48	0.33	0.39	2227
world	0.55	0.62	0.59	4321
accuracy			0.56	22766
macro avg	0.55	0.46	0.49	22766
weighted avg	0.55	0.56	0.55	22766

Confusion Matrix:

```
[46]: array([[ 219,   11,   10,    0,    5,    4,  249,   18,   14,   77],
            [   6, 2802,   12,   13,   27,   43,  677,   83,   63,  300],
            [   1,   34,  672,    6,   10,    1,  224,   55,   48,  167],
            [   0,  158,    7,   45,    9,    5,   96,   51,   39,  168],
            [   2,   66,   21,    6,  294,    2,  147,   48,   73,  171],
            [   1,   46,    3,    0,    3,  746,  267,    8,  109,   20],
            [  84,  522,   74,    1,   40,  178, 3884,  195,  282,  671],
            [   1,  200,   49,   31,   26,   27,  515,  613,   78,  285],
```

```
[ 16, 119, 68, 16, 36, 119, 751, 59, 746, 297],
[ 32, 286, 79, 30, 36, 28, 901, 136, 115, 2678]])
```

Our random forest model does seem to do much better than just one decision tree, with an increase of accuracy from ~45% to ~56%. Additionally, the individual genres' precisions for the most part seemed to get better. This time, classical and metal genres had the highest precision, which, when thinking back to our EDA where the two variables seemed to stand out, makes sense. Similar with decision trees, the model was not as good at classifying hip-hop tracks.

Looking again at feature importance for our random forest:

```
[ ]: # permutation importance
perm_importance_rf = permutation_importance(
    rf_classifier,
    X_test,
    Y_test,
    n_repeats=20,
    random_state=42,
)
```

```
[44]: perm_importance_rf_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance Mean': perm_importance_rf.importances_mean,
    'Importance Std': perm_importance_rf.importances_std
})

perm_importance_rf_df
```

```
[44]:
```

	Feature	Importance Mean	Importance Std
0	popularity	0.119450	0.002112
1	duration_ms	0.034292	0.001582
2	explicit	0.002846	0.000426
3	danceability	0.053180	0.001849
4	energy	0.030636	0.001262
5	key	-0.000356	0.000504
6	loudness	0.016490	0.001054
7	mode	0.004507	0.000626
8	speechiness	0.027034	0.001284
9	acousticness	0.079966	0.001813
10	instrumentalness	0.055049	0.001748
11	liveness	0.004076	0.000793
12	valence	0.031598	0.001160
13	tempo	0.014998	0.001157
14	time_signature	0.000283	0.000386

With our random forest, the importance of specific variables changed slightly from our decision tree, but the most important features, popularity, danceability, and acousticness, stayed about the same.



We also looked into training a KNN model for classification as well, to see how it may do against our decision tree and random forest.

```
[132]: # KNN
knn_classifier = KNeighborsClassifier(n_neighbors=5)
knn_classifier.fit(X_train, Y_train)

# Make predictions
Y_pred_knn = knn_classifier.predict(X_test)

# Evaluate
knn_accuracy = accuracy_score(Y_test, Y_pred_knn)
knn_class_report = classification_report(Y_test, Y_pred_knn, zero_division=1)
knn_conf_matrix = confusion_matrix(Y_test, Y_pred_knn)

print(f'Accuracy: {knn_accuracy}')
print(f'Classification Report: {knn_class_report}')
print(f'Confusion Matrix:')
knn_conf_matrix
```

Accuracy: 0.29438636563296144

Classification Report:                      precision      recall    f1-score      support

classical	0.09	0.10	0.10	607
electronic	0.31	0.47	0.37	4026
folk	0.25	0.24	0.25	1218
hip-hop	0.11	0.06	0.08	578
jazz	0.41	0.29	0.34	830
metal	0.13	0.08	0.10	1203
misc	0.32	0.39	0.35	5931
pop	0.19	0.11	0.14	1825
rock	0.27	0.17	0.21	2227
world	0.33	0.28	0.31	4321
accuracy			0.29	22766
macro avg	0.24	0.22	0.22	22766
weighted avg	0.28	0.29	0.28	22766

Confusion Matrix:

```
[132]: array([[ 59, 143, 29, 9, 6, 26, 221, 14, 24, 76],
 [ 79, 1875, 119, 47, 48, 84, 973, 160, 150, 491],
 [ 33, 252, 290, 21, 26, 32, 338, 33, 45, 148],
 [ 11, 150, 24, 35, 17, 21, 162, 34, 31, 93],
 [ 13, 159, 26, 14, 242, 18, 164, 44, 52, 98],
 [ 41, 268, 44, 22, 16, 95, 416, 46, 87, 168],
 [ 178, 1431, 238, 62, 95, 213, 2309, 248, 346, 811],
 [ 50, 421, 96, 28, 35, 59, 578, 193, 98, 267],
```

```
[ 56, 440, 107, 23, 59, 84, 715, 87, 386, 270],
[ 104, 1004, 173, 53, 43, 115, 1274, 139, 198, 1218]])
```

Unfortunately, our KNN model seems to be a lot worse than both our decision tree and random forest, with an accuracy of only ~29%. Classical, hip-hop, and metal do the worst on this, likely because of the class (genre) imbalance we have from our groupings. KNN doesn't seem to be a very good model for our data.

## 1.7 PCA/Clustering

We tried to conduct PCA and clustering on our data to see if dimensionality reduction or investigation into underlying structures like clusters may relate or help us investigate genre. We tried dimensionality reduction with PCA first.

```
[155]: # data pre-processing
response_variable = 'track_genre'
categorical_columns = [x for x in categorical_columns if x != response_variable]

# drop string columns and response variable
X_train = ohe_grouped_column_transformer_wo_genre.
    ↪transform(grouped_spotify_train)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
Y_train = grouped_spotify_train[response_variable]
```

```
[156]: # do PCA on X
n_components = 10
pca = PCA(n_components=n_components)
X_pca = pca.fit_transform(X_train)
X_pca_df = pd.DataFrame(X_pca, columns=[f'PC{i}' for i in range(n_components)])

explained_variance = pca.explained_variance_ratio_

# Check explained variance ratio for each component
explained_variance = pca.explained_variance_ratio_
print(f"Explained variance ratio for each component: \n {explained_variance}_
    ↪\n")
print(f"Cumulative explained variance: \n {np.cumsum(explained_variance)}")


plt.figure(figsize=(10, 7))
plt.plot(range(1, len(explained_variance) + 1), explained_variance, marker='o',
    ↪linestyle='--')
plt.title('Scree Plot')
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.xticks(range(1, len(explained_variance) + 1))
plt.show()
```

Explained variance ratio for each component:

```
[0.10836837 0.09131586 0.06138773 0.05869306 0.04588949 0.0434164
 0.0361926  0.03418232 0.03292623 0.03291811]
```

Cumulative explained variance:

```
[0.10836837 0.19968423 0.26107196 0.31976503 0.36565451 0.40907091
 0.44526351 0.47944583 0.51237206 0.54529017]
```



Based on the Scree Plot, it seems like taking the first 3 or 4 principal components would be best for still covering a good amount of variance (~0.18 cumulative) and also keeping it simple for easier visualizations. We plotted the first two PCs and colored them by genre to see if there were any interesting relationships.

```
[157]: label_encoder = LabelEncoder()
Y_encoded = label_encoder.fit_transform(Y_train)

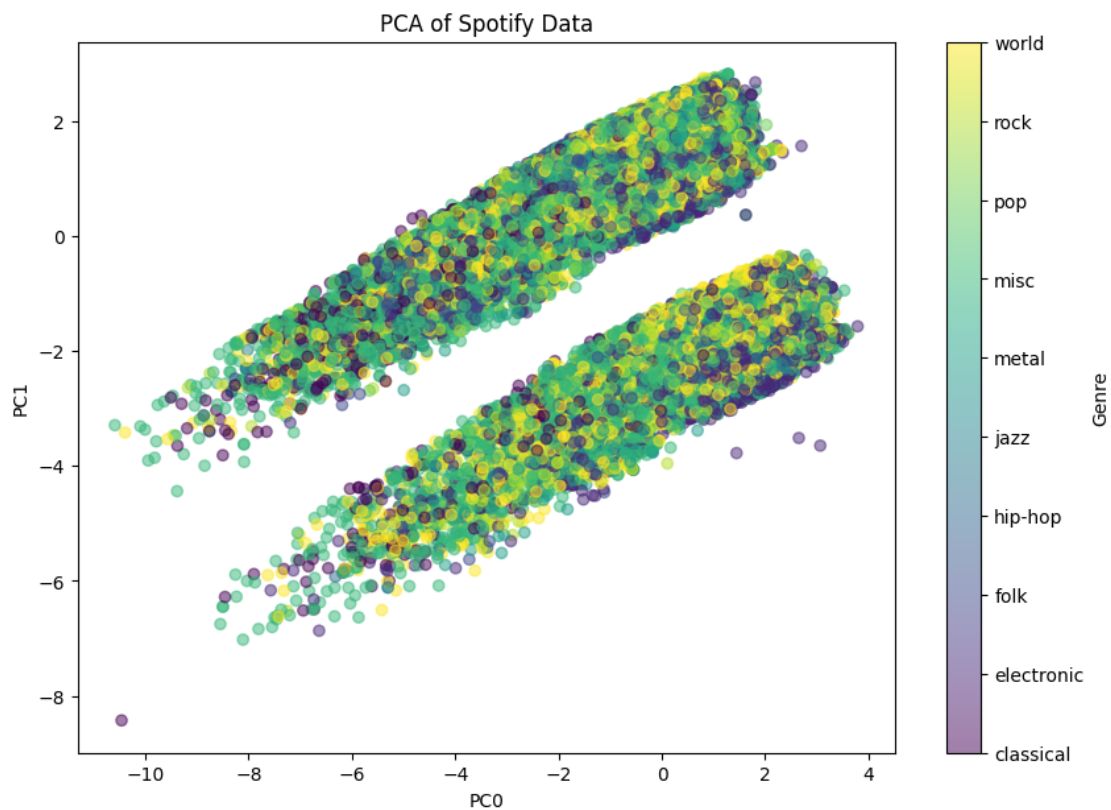
# plot the first two principal components
plt.figure(figsize=(10, 7))
scatter = plt.scatter(
    X_pca_df['PC0'],
    X_pca_df['PC1'],
    c=Y_encoded,
```

```

    cmap='viridis',
    alpha=0.5
)
# label by Y_encoded
cbar = plt.colorbar(scatter)
cbar.set_label('Genre')
cbar.set_ticks(range(len(label_encoder.classes_)))
cbar.set_ticklabels(label_encoder.classes_)

plt.xlabel('PC0')
plt.ylabel('PC1')
plt.title('PCA of Spotify Data')
plt.show()

```



The points do seem like they're getting clustered and colored in some way on the PCA graph which is interesting, but it also looks like there are a lot of overlapping dots, so we may be losing some important dimensions.

We decided to do a regression on the PCA transformed data to see if the PCA improved the fit at all.

```
[159]: # do a regression using PCA transformed data
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score


X_test = ohe_grouped_column_transformer_wo_genre.transform(grouped_spotify_test)
X_test = scaler.transform(X_test)
X_test = pca.transform(X_test)
X_test = pd.DataFrame(X_test, columns=[f'PC{i}' for i in range(n_components)])
Y_test = grouped_spotify_test[response_variable]

model = LogisticRegression(max_iter=1000)
model.fit(X_pca_df, Y_train)
Y_pred = model.predict(X_test)
Y_pred_encoded = label_encoder.transform(Y_pred)

accuracy = accuracy_score(Y_test, Y_pred)
print(f'Accuracy: {accuracy}')

# plot it
plt.figure(figsize=(10, 7))
scatter = plt.scatter(
    X_test['PC0'],
    X_test['PC1'],
    c=Y_pred_encoded,
    cmap='viridis',
    alpha=0.5
)
# label by Y_encoded
cbar = plt.colorbar(scatter)
cbar.set_label('Genre')
cbar.set_ticks(range(len(label_encoder.classes_)))
cbar.set_ticklabels(label_encoder.classes_)
plt.xlabel('PC0')
plt.ylabel('PC1')
plt.title('PCA of Spotify Data')
plt.show()
```

Accuracy: 0.3139769832205921




Though the accuracy isn't very high with our logistic regression fit, the graph reminded us of what Gaussian Mixture Models tend to look like, so we attempted applying a GMM to the PCA transformed data as well.

```
[161]: # gaussian mixture model
n_components = 10
gmm = GaussianMixture(n_components=n_components)
gmm.fit(X_pca_df)
Y_pred = gmm.predict(X_pca_df)

plt.figure(figsize=(10, 7))
scatter = plt.scatter(
    X_pca_df['PC0'],
    X_pca_df['PC1'],
    c=Y_pred,
    cmap='viridis',
    alpha=0.5
)
cbar = plt.colorbar(scatter)
cbar.set_label('Cluster')
cbar.set_ticks(range(n_components))
```

```
cbar.set_ticklabels(range(n_components))

plt.xlabel('PC0')
plt.ylabel('PC1')
plt.title('Gaussian Mixture Model of Spotify Data')
plt.show()
```



```
[162]: # evaluate the model
sil_score = silhouette_score(X_pca_df, Y_pred)
print(sil_score)
```

0.18532988287647995

While the GMM has a pretty nice graph that almost looks right, the silhouette score is not good at all. We decided to try it again with just two genres, pop and not pop (other) to see if it might improve the score.

```
[163]: # Trying the data but with 2 classes: pop and other
Y_binary = Y_train.apply(lambda x: 'pop' if x == 'pop' else 'misc')
```

```


# do GMM
n_components = 2
gmm = GaussianMixture(n_components=n_components)
gmm.fit(X_pca_df)
Y_pred = gmm.predict(X_pca_df)

plt.figure(figsize=(10, 7))
scatter = plt.scatter(
    X_pca_df['PC0'],
    X_pca_df['PC1'],
    c=Y_pred,
    cmap='viridis',
    alpha=0.5
)
cbar = plt.colorbar(scatter)
cbar.set_label('Cluster')
cbar.set_ticks(range(n_components))
cbar.set_ticklabels(range(n_components))

plt.xlabel('PC0')
plt.ylabel('PC1')
plt.title('Gaussian Mixture Model of Spotify Data')
plt.show()

```





```
[164]: # Evaluate
Y_binary_encoded = label_encoder.transform(Y_binary).reshape(-1, 1)
sil_score = silhouette_score(Y_binary_encoded, Y_pred)

print(sil_score)
```

-0.030537843090037743

Unfortunately, our silhouette score is still less than zero, indicating that our clustering with the GMM is not very good. We decided to try other methods of clustering, K-means and agglomerative, as well to see if they would be better.

Note that when working with the original data that has 114 genres, it's very evenly distributed (1000 samples per genre). This may imply that the sampling was not random, and was stratified via the genre subpopulations. This is likely not proportional to the actual song population, which may lead to bias since underrepresented genres are now equally represented with overrepresented genres. Since we have no way to recover original proportions without using external data, we will simply have to be weary of the results.

Our goal with clustering is to be able to see if it might be able to help us in classifying or predicting the genre of the song. We first start to see if there is any obvious clustering using both agglomerative

(hierarchical) and k-means clustering.

```
[175]: X_train = ohe_grouped_column_transformer_wo_genre.  
        ↪transform(grouped_spotify_train)  
X_test = ohe_grouped_column_transformer_wo_genre.transform(grouped_spotify_test)  
  
scaler = StandardScaler()  
X_train = pd.DataFrame(  
    scaler.fit_transform(X_train),  
    index=X_train.index,  
    columns=X_train.columns  
)  
X_test = pd.DataFrame(  
    scaler.transform(X_test),  
    index=X_test.index,  
    columns=X_test.columns  
)  
  
Y_train = grouped_spotify_train[response_variable]  
Y_test = grouped_spotify_test[response_variable]
```

```
[172]: # sample the training set to speed up hierarchical clustering  
sample_size = 0.3  
hierarchical_sample_split = StratifiedShuffleSplit(  
    n_splits=1,  
    test_size=1 - sample_size,  
    random_state=42  
)  
gmm_sample_size = 0.8  
  
for train_index, _ in hierarchical_sample_split.split(X_train, Y_train):  
    X_sample = X_train.iloc[train_index].copy()  
    Y_sample = Y_train.iloc[train_index].copy()
```

```
[173]: n_clusters = 10  
  
kmeans = KMeans(  
    n_clusters=n_clusters,  
    random_state=0  
)  
hierarchical = AgglomerativeClustering(  
    n_clusters=n_clusters,  
    metric='euclidean',  
    linkage='ward'  
)  
gmm = GaussianMixture(  
    n_components=n_clusters,
```

```

        random_state=0
    )

    hierarchical.fit(X_sample)
    kmeans.fit(X_sample)
    gmm.fit(X_train)

```

[173]: GaussianMixture(n\_components=10, random\_state=0)

```

[176]: X_sample["hcluster"] = hierarchical.labels_
X_sample["kcluster"] = kmeans.labels_
Y_pred_gmm = gmm.predict(X_test)

# encode Y_test_gmm labels using LabelEncoder
label_encoder = LabelEncoder()
label_encoder.fit(Y_train)
Y_train_encoded = label_encoder.transform(Y_train)

```

```

[ ]: # Analyze
hclust_kclust_ari = adjusted_rand_score(X_sample["hcluster"],
    ↪X_sample["kcluster"])

hclust_y_ari = adjusted_rand_score(X_sample["hcluster"], Y_sample)
kclust_y_ari = adjusted_rand_score(X_sample["kcluster"], Y_sample)
kclust_y_ari = adjusted_rand_score(Y_test, Y_pred_gmm)

hclust_silhouette = silhouette_score(X_sample, X_sample["hcluster"])
kclust_silhouette = silhouette_score(X_sample, X_sample["kcluster"])
gmm_silhouette = silhouette_score(X_test, Y_pred_gmm)

print(f'''
Hierarchical clustering vs KMeans clustering ARI: {hclust_kclust_ari}

Hierarchical clustering vs true labels ARI: {hclust_y_ari}
KMeans clustering vs true labels ARI: {kclust_y_ari}
GMM clustering vs true labels ARI: {kclust_y_ari}

Hierarchical clustering silhouette score: {hclust_silhouette}
KMeans clustering silhouette score: {kclust_silhouette}
GMM clustering silhouette score: {gmm_silhouette}
''')

```

Hierarchical clustering vs KMeans clustering ARI: 0.4500737924478871

Hierarchical clustering vs true labels ARI: 0.0019950842701675522

KMeans clustering vs true labels ARI: 0.0022420916160579103

GMM clustering vs true labels ARI: 0.0022420916160579103

Hierarchical clustering silhouette score: 0.16310140621490865

KMeans clustering silhouette score: 0.22201717965090853

GMM clustering silhouette score: 0.1694012135651898

```
[180]: pd.set_option('display.max_columns', None)
X_sample.groupby("hcluster").count()
```

```
[180]:
```

	key_0	key_1	key_2	key_3	key_4	key_5	key_6	key_7	key_8	\
hcluster										
0	4098	4098	4098	4098	4098	4098	4098	4098	4098	
1	4902	4902	4902	4902	4902	4902	4902	4902	4902	
2	4188	4188	4188	4188	4188	4188	4188	4188	4188	
3	1190	1190	1190	1190	1190	1190	1190	1190	1190	
4	1983	1983	1983	1983	1983	1983	1983	1983	1983	
5	1515	1515	1515	1515	1515	1515	1515	1515	1515	
6	1497	1497	1497	1497	1497	1497	1497	1497	1497	
7	334	334	334	334	334	334	334	334	334	
8	614	614	614	614	614	614	614	614	614	
9	167	167	167	167	167	167	167	167	167	

	key_9	key_10	key_11	mode_0	mode_1	time_signature_0	\
hcluster							
0	4098	4098	4098	4098	4098	4098	
1	4902	4902	4902	4902	4902	4902	
2	4188	4188	4188	4188	4188	4188	
3	1190	1190	1190	1190	1190	1190	
4	1983	1983	1983	1983	1983	1983	
5	1515	1515	1515	1515	1515	1515	
6	1497	1497	1497	1497	1497	1497	
7	334	334	334	334	334	334	
8	614	614	614	614	614	614	
9	167	167	167	167	167	167	

	time_signature_1	time_signature_3	time_signature_4	\
hcluster				
0	4098	4098	4098	
1	4902	4902	4902	
2	4188	4188	4188	
3	1190	1190	1190	
4	1983	1983	1983	
5	1515	1515	1515	
6	1497	1497	1497	
7	334	334	334	
8	614	614	614	
9	167	167	167	

	time_signature_5	popularity	duration_ms	explicit	danceability	\
hcluster						
0	4098	4098	4098	4098	4098	
1	4902	4902	4902	4902	4902	
2	4188	4188	4188	4188	4188	
3	1190	1190	1190	1190	1190	
4	1983	1983	1983	1983	1983	
5	1515	1515	1515	1515	1515	
6	1497	1497	1497	1497	1497	
7	334	334	334	334	334	
8	614	614	614	614	614	
9	167	167	167	167	167	

	energy	key	loudness	mode	speechiness	acousticness	\
hcluster							
0	4098	4098	4098	4098	4098	4098	
1	4902	4902	4902	4902	4902	4902	
2	4188	4188	4188	4188	4188	4188	
3	1190	1190	1190	1190	1190	1190	
4	1983	1983	1983	1983	1983	1983	
5	1515	1515	1515	1515	1515	1515	
6	1497	1497	1497	1497	1497	1497	
7	334	334	334	334	334	334	
8	614	614	614	614	614	614	
9	167	167	167	167	167	167	

	instrumentalness	liveness	valence	tempo	time_signature	kcluster
hcluster						
0	4098	4098	4098	4098	4098	4098
1	4902	4902	4902	4902	4902	4902
2	4188	4188	4188	4188	4188	4188
3	1190	1190	1190	1190	1190	1190
4	1983	1983	1983	1983	1983	1983
5	1515	1515	1515	1515	1515	1515
6	1497	1497	1497	1497	1497	1497
7	334	334	334	334	334	334
8	614	614	614	614	614	614
9	167	167	167	167	167	167

Based on our investigations, the K-means and Agglomerative clustering did ok, with silhouette scores for both methods being ~0.5. However, like before, GMM did terribly. However, the adjusted rand index when compared with track genre was negligible, meaning that the clusters formed do not at all correspond to genre, so clustering will probably not help us with our classification problem.

## 1.8 Neural Networks

We then applied neural networks to see if they could approximate the genre as a function of the other features.

We first put one-hot-encoded the data, scaled it, and put it onto the GPU (if available). We used a batch size of 64 for data loading.

```
[185]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

/home/ketexon/programming/csm148-spotiflies/.venv/lib/python3.12/site-
packages/torch/cuda/__init__.py:716: UserWarning: Can't initialize NVML
  warnings.warn("Can't initialize NVML")
/home/ketexon/programming/csm148-spotiflies/.venv/lib/python3.12/site-
packages/torch/cuda/__init__.py:905: UserWarning: CUDA initialization:
Unexpected error from cudaGetDeviceCount(). Did you run some cuda functions
before calling NumCudaDevices() that might have already set an error? Error 803:
system has unsupported display driver / cuda driver combination (Triggered
internally at ../c10/cuda/CUDAFunctions.cpp:108.)
  r = torch._C._cuda_getDeviceCount() if nvml_count < 0 else nvml_count

[186]: response_variable = "track_genre"

# remove string columns and one-hot encode
Y_train = grouped_spotify_train[response_variable]
Y_val = grouped_spotify_val[response_variable]

X_train = ohe_grouped_column_transformer_wo_genre.
    ↪transform(grouped_spotify_train)
X_val = ohe_grouped_column_transformer_wo_genre.transform(grouped_spotify_val)

label_encoder = LabelEncoder()
label_encoder.fit(grouped_spotify[response_variable])
Y_train = label_encoder.transform(Y_train)
Y_val = label_encoder.transform(Y_val)

# create class weights
Y_train_unique = np.unique(Y_train, return_counts=True)
class_weights = len(Y_train) / (Y_train_unique[1] * len(Y_train_unique[1]))
for i, w in enumerate(class_weights):
    label_encoded = Y_train_unique[0][i]
    label_decoded = label_encoder.inverse_transform([label_encoded])[0]
    w_str = f"{w:.3f}"
    count_str = f"{Y_train_unique[1][i]:,}"
    print(f"{label_decoded + ':': <16}{w_str: <16}{count_str}")

# scale
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
```

```

X_val = scaler.transform(X_val)

# convert to tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32).to(device)
X_val_tensor = torch.tensor(X_val, dtype=torch.float32).to(device)
Y_train_tensor = torch.tensor(Y_train, dtype=torch.long).to(device)
Y_val_tensor = torch.tensor(Y_val, dtype=torch.long).to(device)

class_weights_tensor = torch.tensor(class_weights, dtype=torch.float32).
    ↪to(device)

# create dataset
train_dataset = TensorDataset(X_train_tensor, Y_train_tensor)
validation_dataset = TensorDataset(X_val_tensor, Y_val_tensor)

# create dataloaders
batch_size = 64
num_workers = 12 if is_linux else 0
train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=num_workers,
)
validation_loader = DataLoader(
    validation_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=num_workers,
)

```

classical:	3.820	1,788
electronic:	0.569	12,002
folk:	1.895	3,604
hip-hop:	3.744	1,824
jazz:	2.854	2,393
metal:	1.876	3,640
misc:	0.380	17,962
pop:	1.268	5,386
rock:	1.040	6,569
world:	0.520	13,128

Then we defined our model. Via trial and error, we found that using ReLU caused slow training (since the gradient is 0 if the value is negative), and including dropout added to much variation to the training loss. In addition, switching from softmax to log softmax helped the early stopping algorithm from stopping too early. Thus, we kept a simple model, with 2 fully connected hidden layers, with sigmoid activation for the hidden layers and log softmax for the output.

We used the Adam optimizer with weight decay for regularization using cross entropy loss. We also used a learning rate scheduler which reduced learning rate by 90% every 10 epochs.

```
[187]: # Create model
class SpotifyModel(pl.LightningModule):
    def __init__(
        self,
        input_dim: int,
        output_dim: int,
        lr: float = 0.01,
        class_weights: torch.tensor = None,
        hidden_dim1: int = 64,
        hidden_dim2: int = 64,
        weight_decay: float = 0.01,
    ):
        super().__init__()
        self.lr = lr
        self.model = nn.Sequential(
            nn.Linear(input_dim, hidden_dim1),
            nn.Sigmoid(),
            nn.Linear(hidden_dim1, hidden_dim2),
            nn.Sigmoid(),
            nn.Linear(hidden_dim2, output_dim),
            nn.LogSoftmax(dim=1),
        )
        self.n_classes = output_dim
        self.weight_decay = weight_decay
        self.epoch_metrics = []
        self.train_metrics_stack = []
        self.validation_metrics_stack = []
        self.loss = nn.CrossEntropyLoss(weight=class_weights)

    def forward(self, x):
        return self.model(x)

    def training_step(self, batch, batch_idx):
        x, y = batch
        y_hat = self(x)
        loss = self.loss(y_hat, y)
        self.log("train_loss", loss)
        self.train_metrics_stack.append({
            'train_loss': loss,
        })
        return loss

    def validation_step(self, batch, batch_idx):
        x, y = batch
```



```

        y_hat = self(x)
        loss = self.loss(y_hat, y)
        preds = torch.argmax(y_hat, dim=1)
        acc = (preds == y).float().mean()
        self.validation_metrics_stack.append({
            'val_loss': loss,
            'val_acc': acc
        })
        self.log("val_loss", loss)
        return loss

    def on_validation_epoch_end(self):
        # sum in validation stack
        if len(self.validation_metrics_stack) == 0:
            return
        if len(self.train_metrics_stack) == 0:
            return
        val_loss = torch.stack([x['val_loss'] for x in self.
↪validation_metrics_stack]).mean()
        val_acc = torch.stack([x['val_acc'] for x in self.
↪validation_metrics_stack]).mean()
        train_loss = torch.stack([x['train_loss'] for x in self.
↪train_metrics_stack]).mean()
        self.epoch_metrics.append({
            'val_loss': val_loss.cpu(),
            'val_acc': val_acc.cpu(),
            'train_loss': train_loss.cpu(),
        })
        # clear stack
        self.validation_metrics_stack = []
        self.train_metrics_stack = []

    def configure_optimizers(self):
        optimizer = optim.Adam(self.parameters(), lr=self.lr, weight_decay=self.
↪weight_decay)
        scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.
↪1)
        return [optimizer], [scheduler]

```

To pick an initial learning rate, we used PyTorch's `lr_find` function, which gave us an initial learning rate of about 0.002. We then trained the model, which stopped early at 27 epochs.

```

[ ]: input_dim = X_train.shape[1]
    output_dim = len(
        label_encoder.classes_
    )

```

```

lr = 0.001

model = SpotifyModel(
    input_dim,
    output_dim,
    lr=lr,
    class_weights=class_weights_tensor
).to(device)

# train model
early_stopping = pl_callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    min_delta=0.0001,
    mode='min'
)

use_checkpoint = False
ckpt_path="epoch=12-step=18499.ckpt" if use_checkpoint else None
trainer = pl.Trainer(
    max_epochs=100,
    callbacks=[
        early_stopping
    ]
)

# tune the model
from lightning.pytorch.tuner.tuning import Tuner
tuner = Tuner(trainer)

lr_finder = tuner.lr_find(
    model,
    train_dataloaders=train_loader,
    val_dataloaders=validation_loader,
    min_lr=1e-10,
    max_lr=1e-1,
    num_training=100,
)

model.lr = lr_finder.suggestion()
print(f"Learning rate: {model.lr}")

trainer.fit(
    model,
    train_loader,
    validation_loader,
    ckpt_path=ckpt_path

```

```
)
```

```
/home/ketexon/programming/csm148-spotiflies/.venv/lib/python3.12/site-
packages/torch/cuda/__init__.py:716: UserWarning: Can't initialize NVML
  warnings.warn("Can't initialize NVML")
GPU available: False, used: False
TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
Finding best initial lr: 96%|      | 96/100 [00:00<00:00,
138.32it/s] `Trainer.fit` stopped: `max_steps=100` reached.
Finding best initial lr: 100%|     | 100/100 [00:00<00:00, 132.55it/s]
Learning rate set to 0.0019498445997580478
Restoring states from the checkpoint path at /home/ketexon/programming/csm148-
spotiflies/.lr_find_71f92498-ffbf-4528-aab4-a052b2bdc14c.ckpt
Restored all states from the checkpoint at /home/ketexon/programming/csm148-
spotiflies/.lr_find_71f92498-ffbf-4528-aab4-a052b2bdc14c.ckpt
```

	Name	Type	Params	Mode
0	model	Sequential	6.7 K	train
1	loss	CrossEntropyLoss	0	train

6.7 K	Trainable params
0	Non-trainable params
6.7 K	Total params
0.027	Total estimated model params size (MB)
8	Modules in train mode
0	Modules in eval mode

```
Learning rate: 0.0019498445997580478
Epoch 27: 100%|      | 1423/1423 [00:14<00:00, 100.88it/s, v_num=55]
```

We then plotted a classification report on the model.

```
[ ]: # evaluate model
model.eval()
with torch.no_grad():
    outputs = model.to(device)(X_val_tensor.to(device)).cpu()
    _, predicted = torch.max(outputs.data, 1)
    value_counts = pd.DataFrame(predicted).value_counts()
    for label in label_encoder.classes_:
        class_encoded = label_encoder.transform([label])[0]
        count = value_counts.get(class_encoded, 0)
        print(f"{label}: {count}")
    class_report = classification_report(
        Y_val,
        predicted,
        labels=range(len(label_encoder.classes_)),
```

```

        target_names=label_encoder.classes_,
        zero_division=1
    )
model.train()

print(class_report)

```

```

classical: 4629
electronic: 4627
folk: 2029
hip-hop: 0
jazz: 890
metal: 6415
misc: 3462
pop: 0
rock: 609
world: 105

```

	precision	recall	f1-score	support
classical	0.10	0.77	0.18	607
electronic	0.26	0.30	0.28	4026
folk	0.09	0.15	0.11	1218
hip-hop	1.00	0.00	0.00	578
jazz	0.07	0.07	0.07	830
metal	0.17	0.90	0.28	1203
misc	0.25	0.15	0.19	5931
pop	1.00	0.00	0.00	1825
rock	0.09	0.03	0.04	2227
world	0.30	0.01	0.01	4321
accuracy			0.17	22766
macro avg	0.33	0.24	0.12	22766
weighted avg	0.30	0.17	0.13	22766

The classification shows pretty bad results, with overall 17% accuracy (which is hardly better than uniform, at 10%).


However, something interesting of note is that the recall for classical and metal was significantly high. This means that if the model predicted that a song was metal or classical, it was likely correct. Something else of note was that pop and hip hop, the most and least represented classes, had 0 predictions. This was stranger, but through other models, we saw that the class imbalance disproportionately made more represented classes be predicted. When we added class weights, this might have overcompensated for pop, leading to pop having poor predictions, and undercompensated for hip-hop.

```


[ ]: # Graph epoch metrics
import matplotlib.pyplot as plt

```

```
# plot epoch to val loss and train loss
plt.plot([x['val_loss'].cpu() for x in model.epoch_metrics], label='val_loss')
plt.plot([x['train_loss'].cpu() for x in model.epoch_metrics],
         label='train_loss')
plt.legend()
plt.show()
```



```
[ ]: # plot accuracy per epoch
plt.plot([x['val_acc'].cpu() for x in model.epoch_metrics], label='val_acc')
plt.legend()
plt.show()
```



Plotting the validation and training loss, as well as the model accuracy over epochs, we can see that the loss did shrink over time, with accuracy never really increasing. This could be an indicator that cross entropy was not a good loss function.